



# Building Scalable Applications with Hazelcast

**FUAD MALIKOV**  
**CO-FOUNDER**

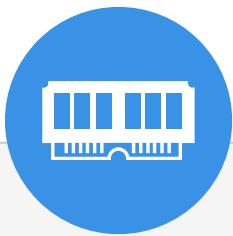


# What Is Hazelcast?

**Hazelcast is a distributed,  
highly available and scalable  
Open Source In-Memory Data Grid**



# In Memory Data Grid



In Memory  
Data **Storage**



In Memory  
Data **Messaging**



In Memory  
Data **Computing**





# java.util.Map

```
import java.util.HashMap;
import java.util.Map;

public static void main(String[] args) {
    Map<Integer, String> map = new HashMap<>();
    map.put(1, "Paris");
    map.put(2, "London");
    map.put(3, "San Francisco");

    String oldValue = map.remove(2);
}
```



# java.util.concurrent.ConcurrentMap

```
import java.util.concurrent.ConcurrentHashMap;
import java.util.concurrent.ConcurrentMap;

public static void main(String[] args) {
    ConcurrentMap<Integer, String> map = new ConcurrentHashMap<>();
    map.put(1, "Paris");
    map.put(2, "London");
    map.put(3, "San Francisco");

    String oldValue = map.remove(2);
}
```



# Distributed Map

```
import java.util.concurrent.ConcurrentMap;
import com.hazelcast.core.Hazelcast;
import com.hazelcast.core.HazelcastInstance;

public static void main(String[] args) {
    HazelcastInstance h = Hazelcast.newHazelcastInstance();

    ConcurrentMap<Integer, String> map = h.getMap("myMap");
    map.put(1, "Paris");
    map.put(2, "London");
    map.put(3, "San Francisco");

    String oldValue = map.remove(2);
}
```



# Ecosystem Traction

*Dozens of Commercial and Open Source Projects Embed Hazelcast*





# Demo



# Why Hazelcast?



**Scale-out Computing** enables cluster capacity to be increased or decreased on-demand



**Resilience** with automatic recovery from member failures without losing data while minimizing performance impact on running applications



**Programming Model** provides a way for developers to easily program a cluster application as if it is a single process



**Fast Application Performance** enables very large data sets to be held in main memory for real-time performance



# Rebalance Data on New Node



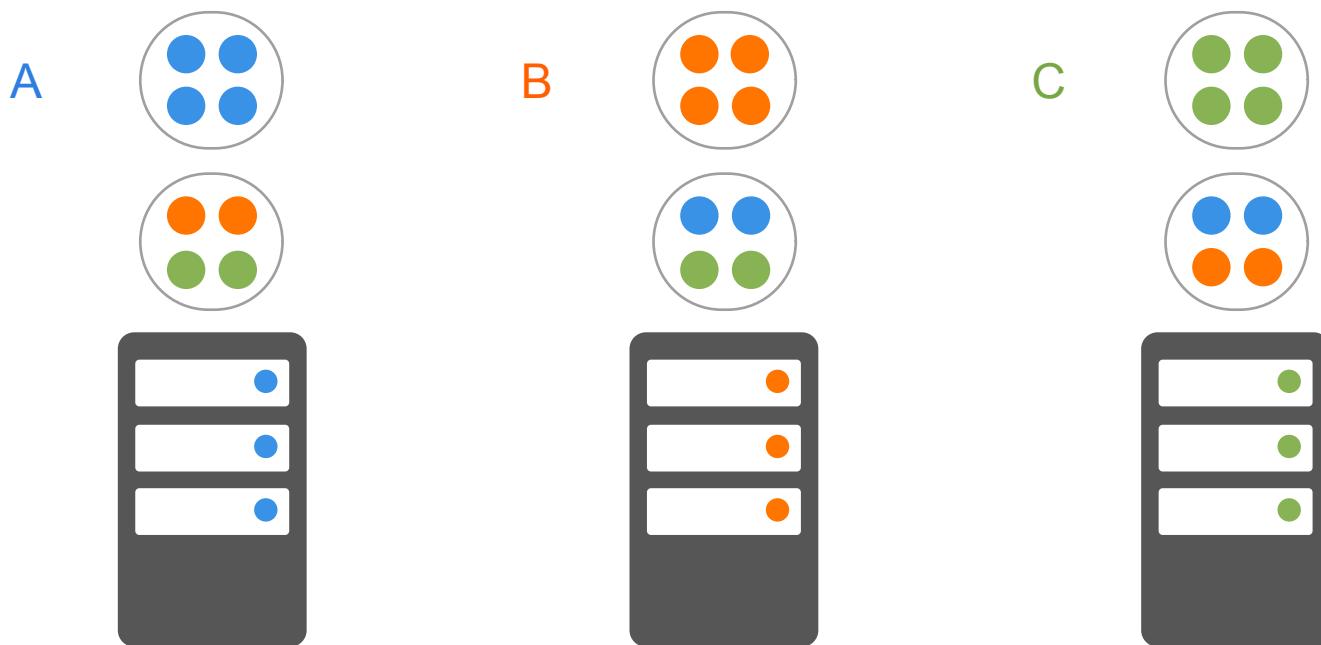
# Distributed Maps

Fixed number of partitions (default 271)

Each key falls into a partition

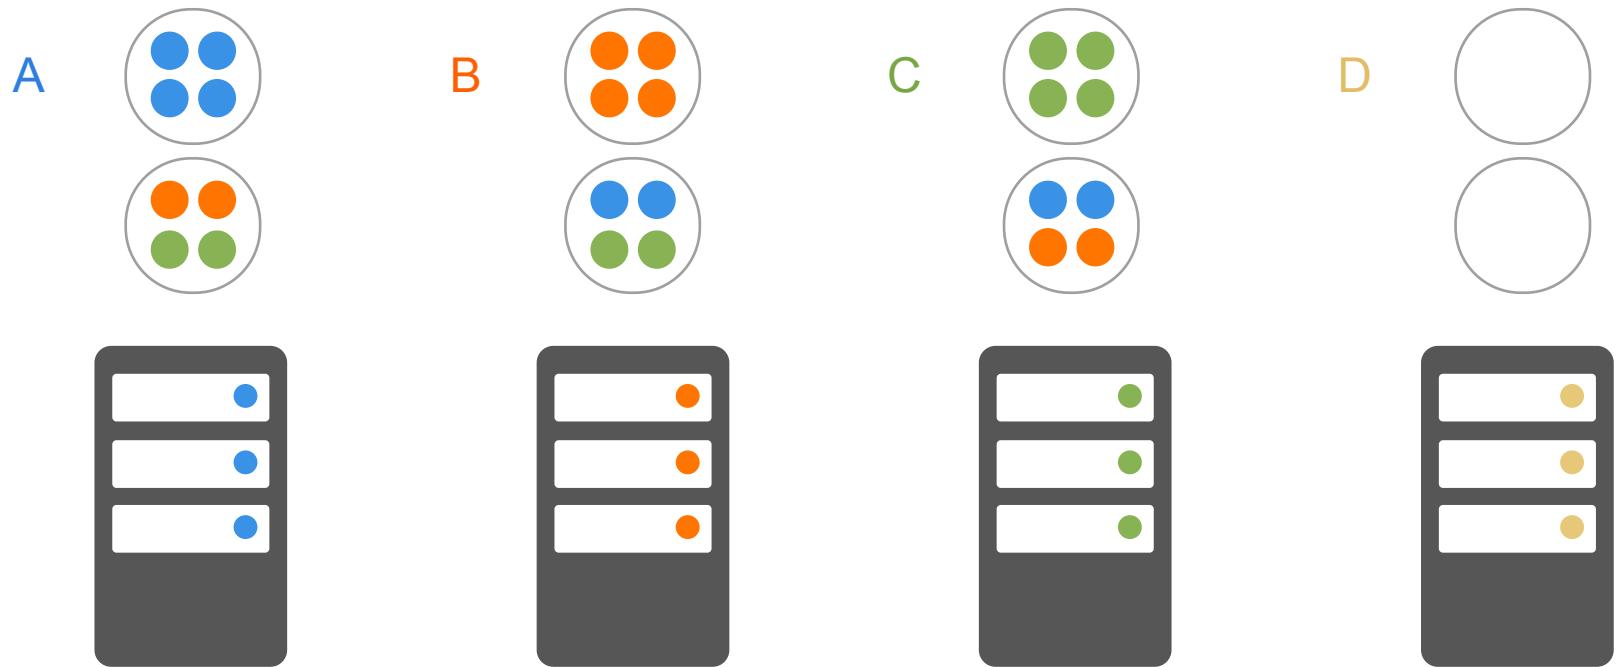
*partitionId = hash(keyData)%PARTITION\_COUNT*

Partition ownerships are reassigned upon membership change



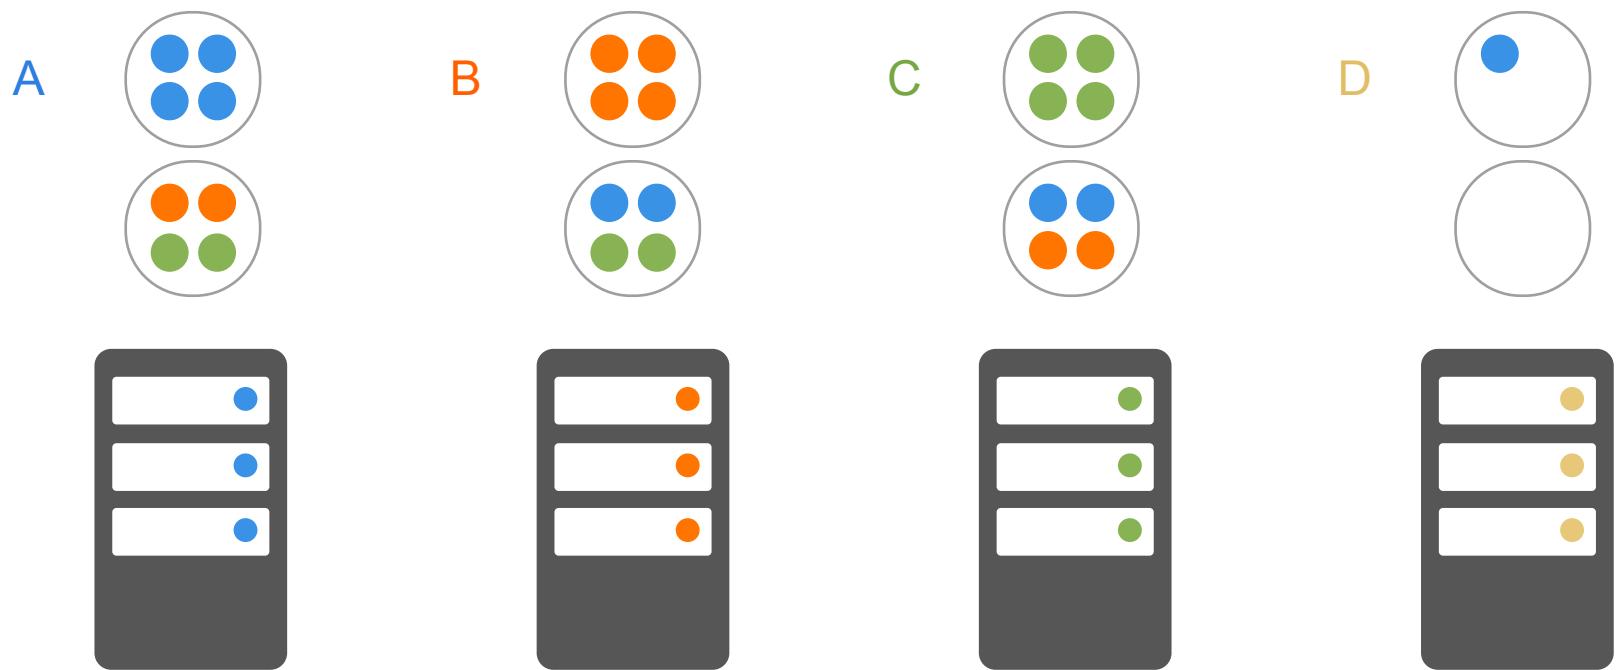


# New Node Added



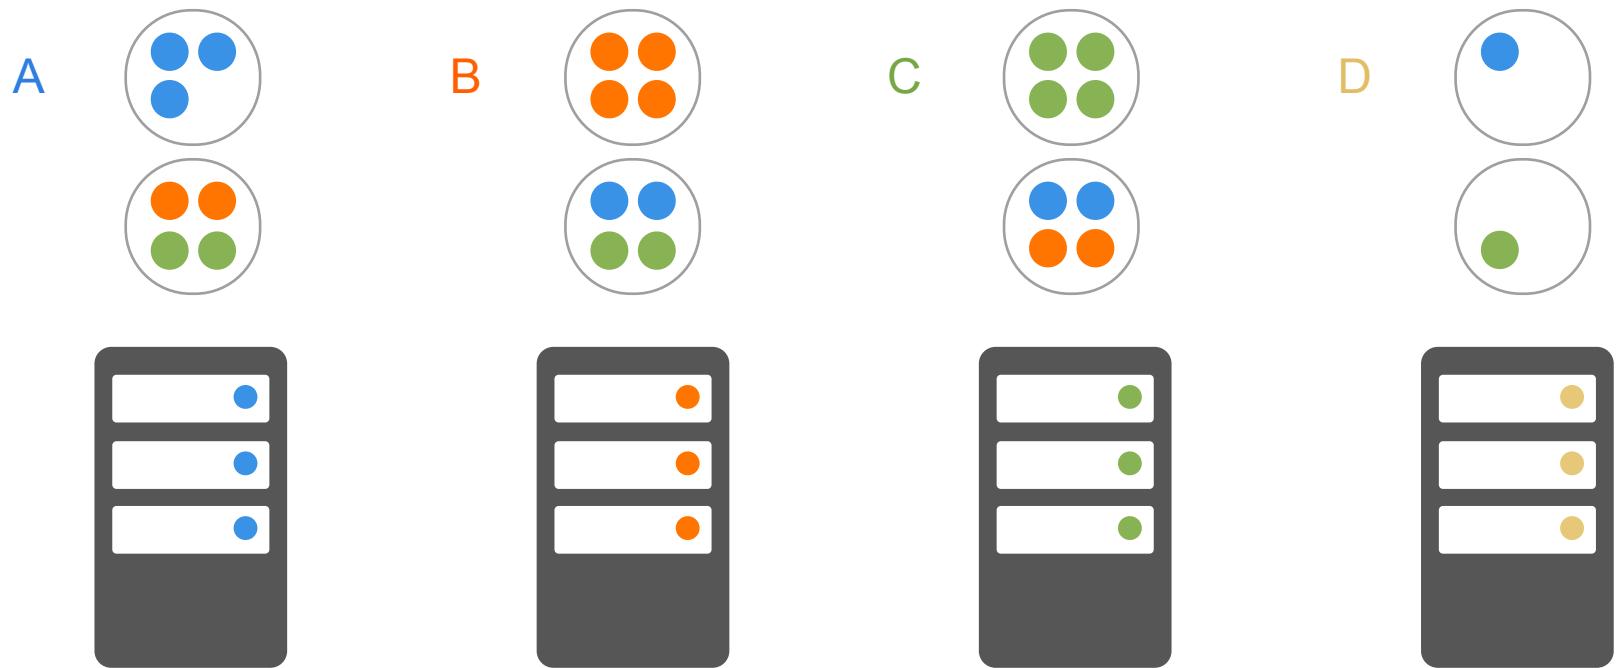


# Migration



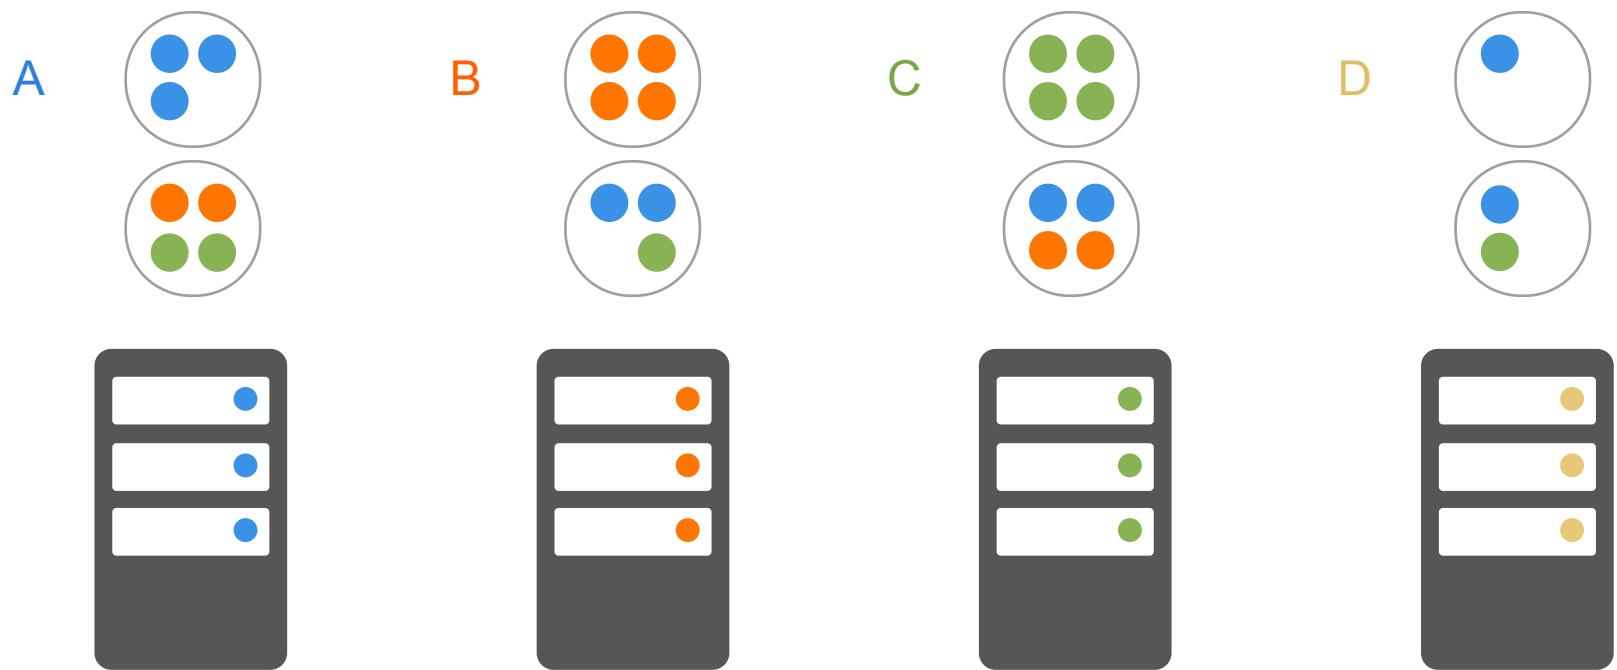


# Migration



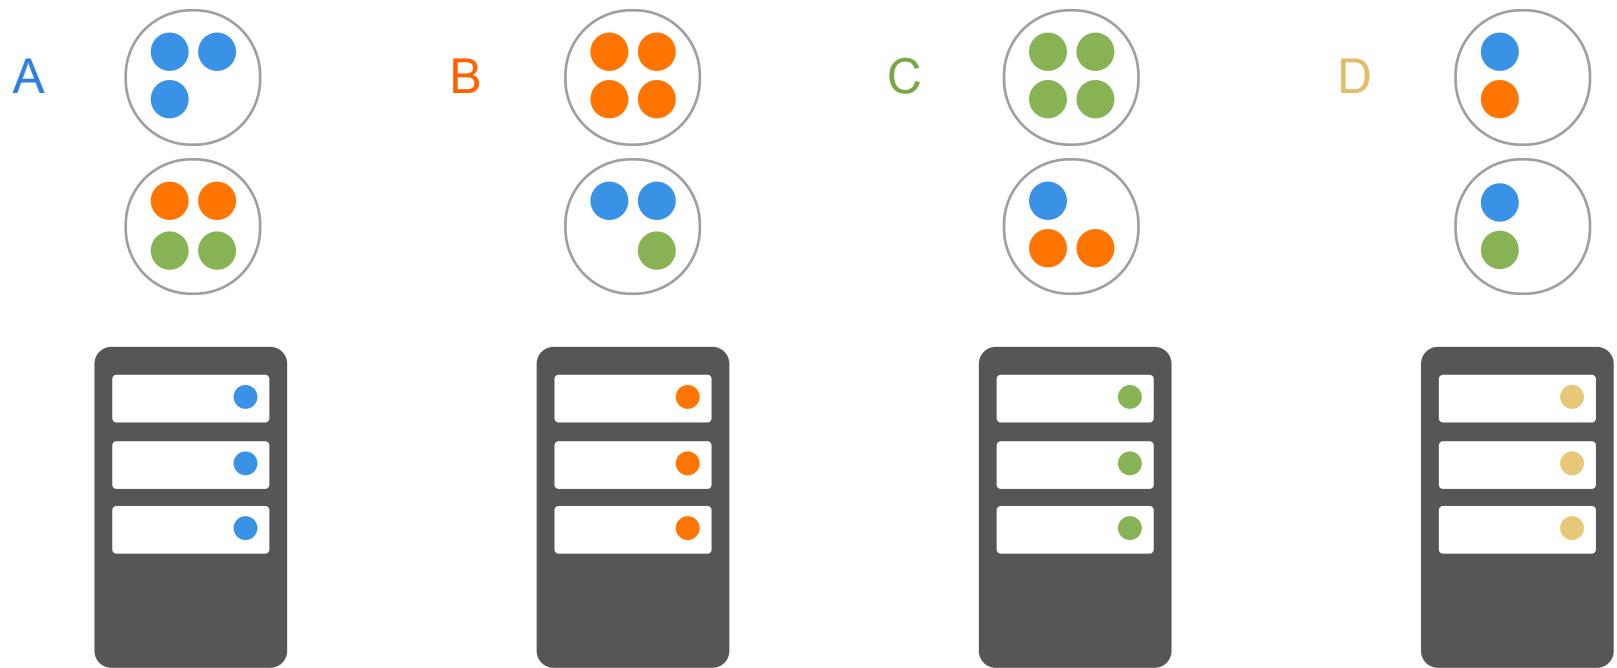


# Migration



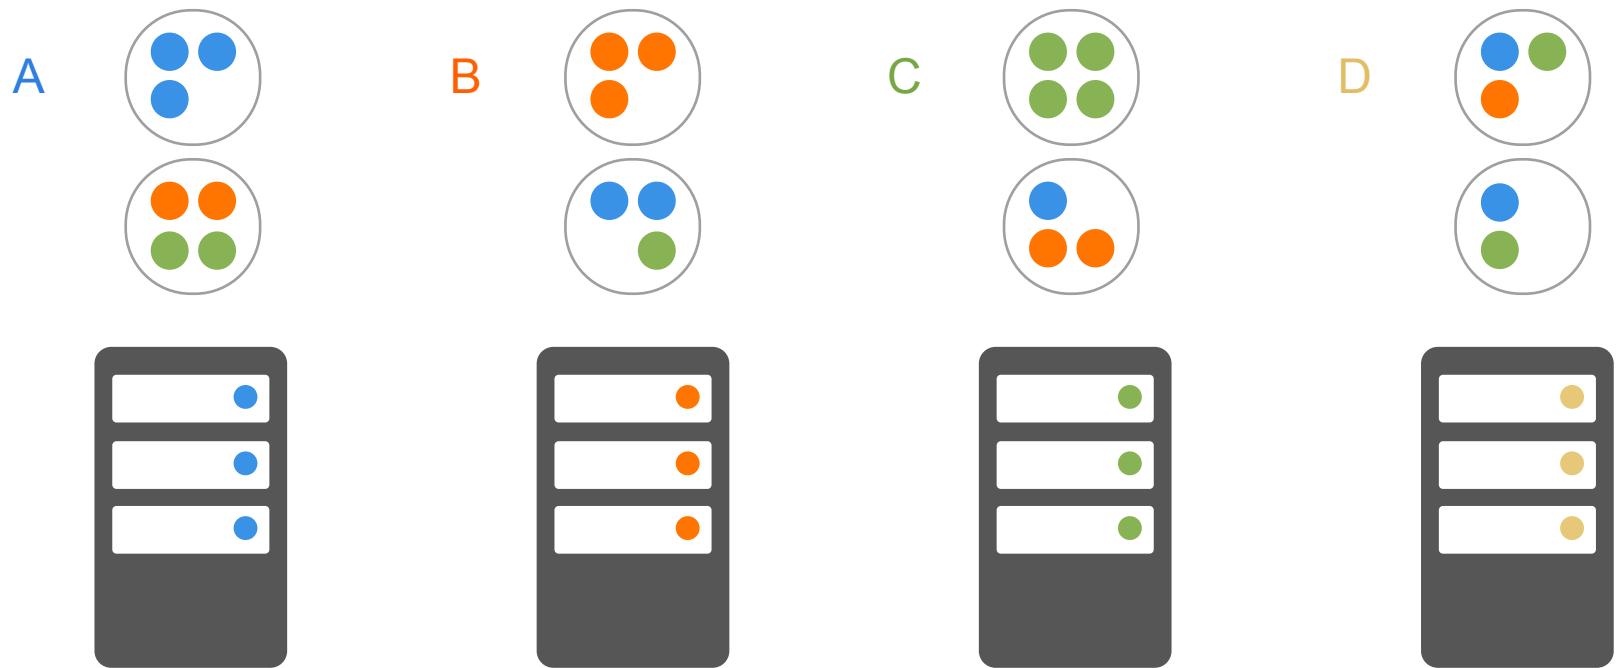


# Migration



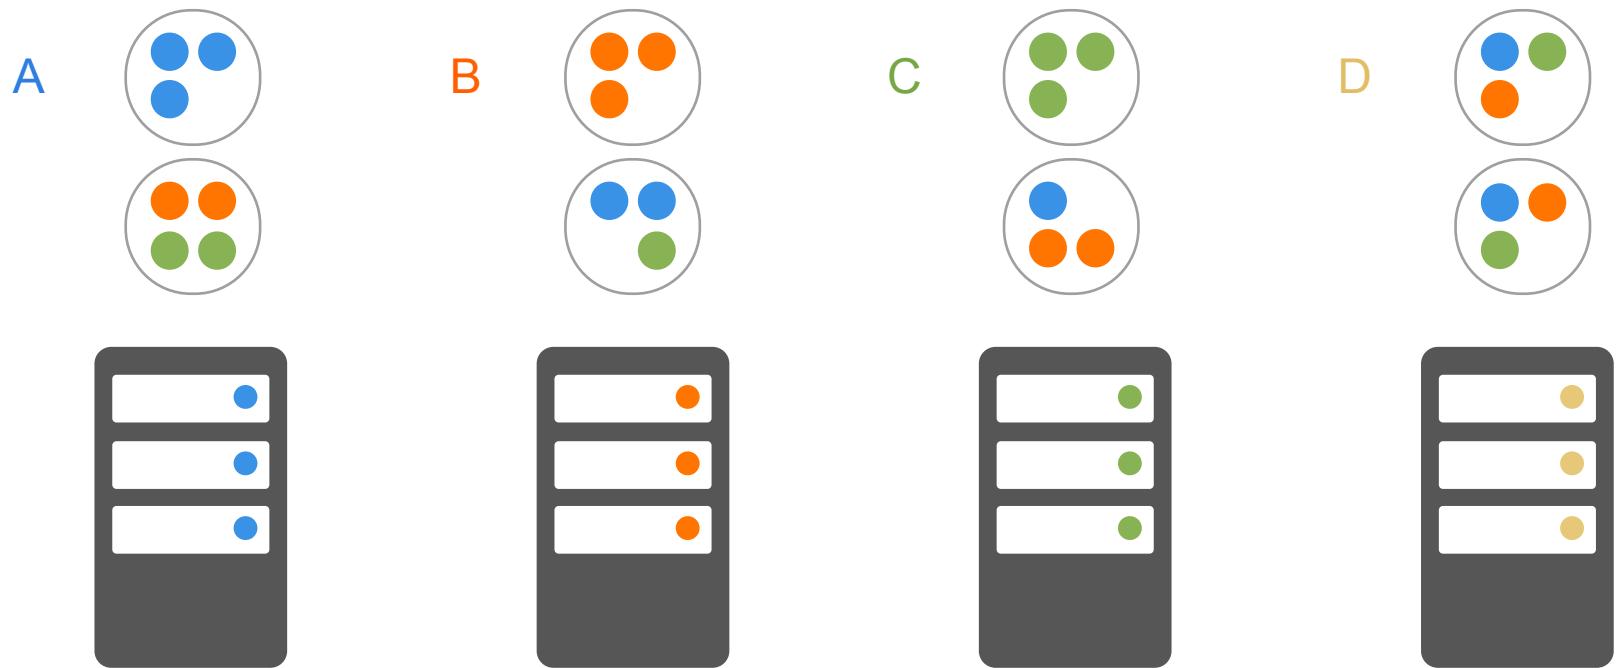


# Migration



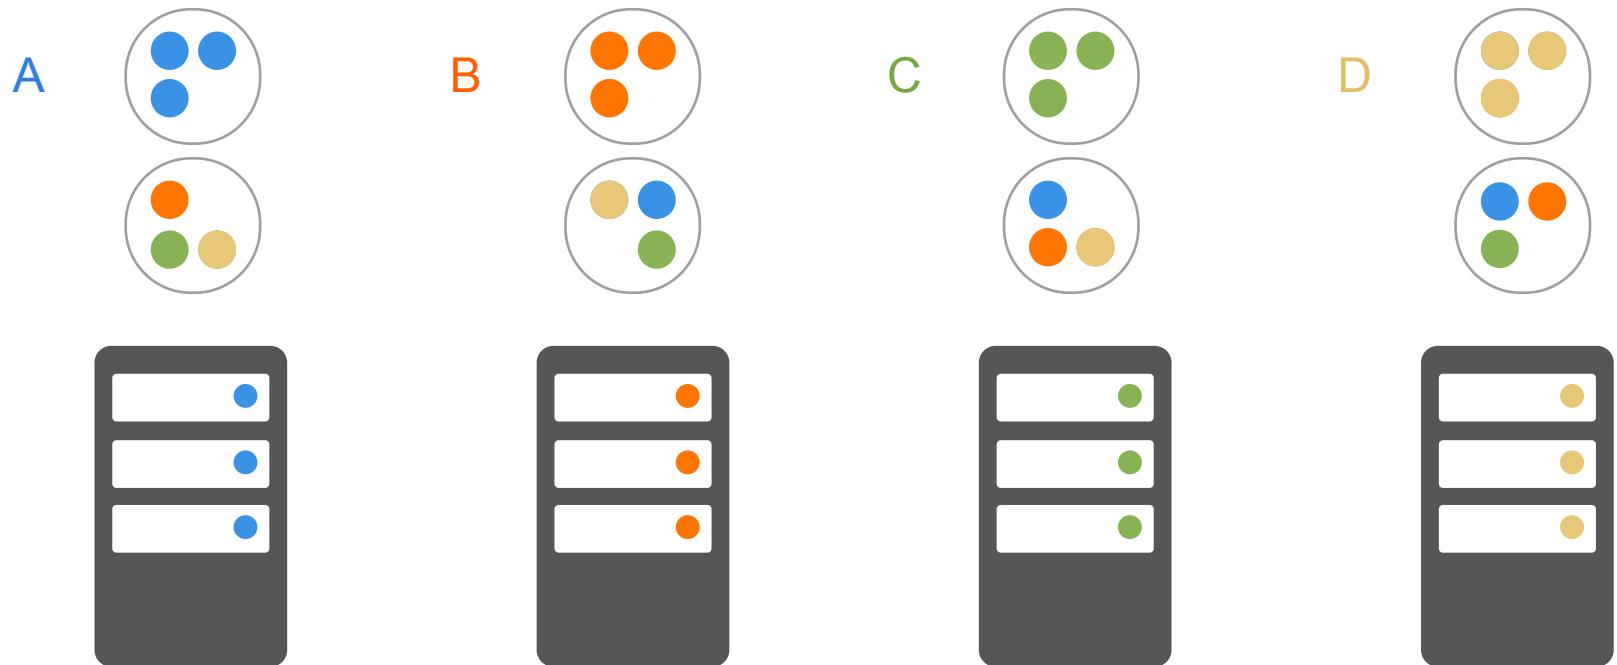


# Migration





# Migration Complete



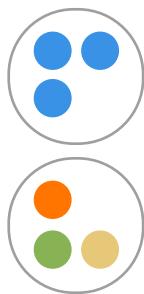


# Data Safety when Node Dies



# Node Crashes

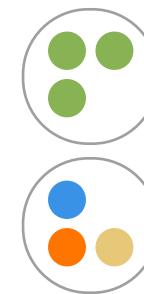
A



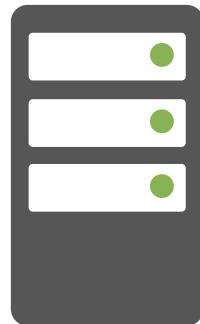
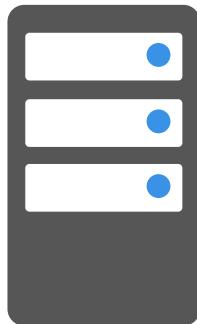
B



C

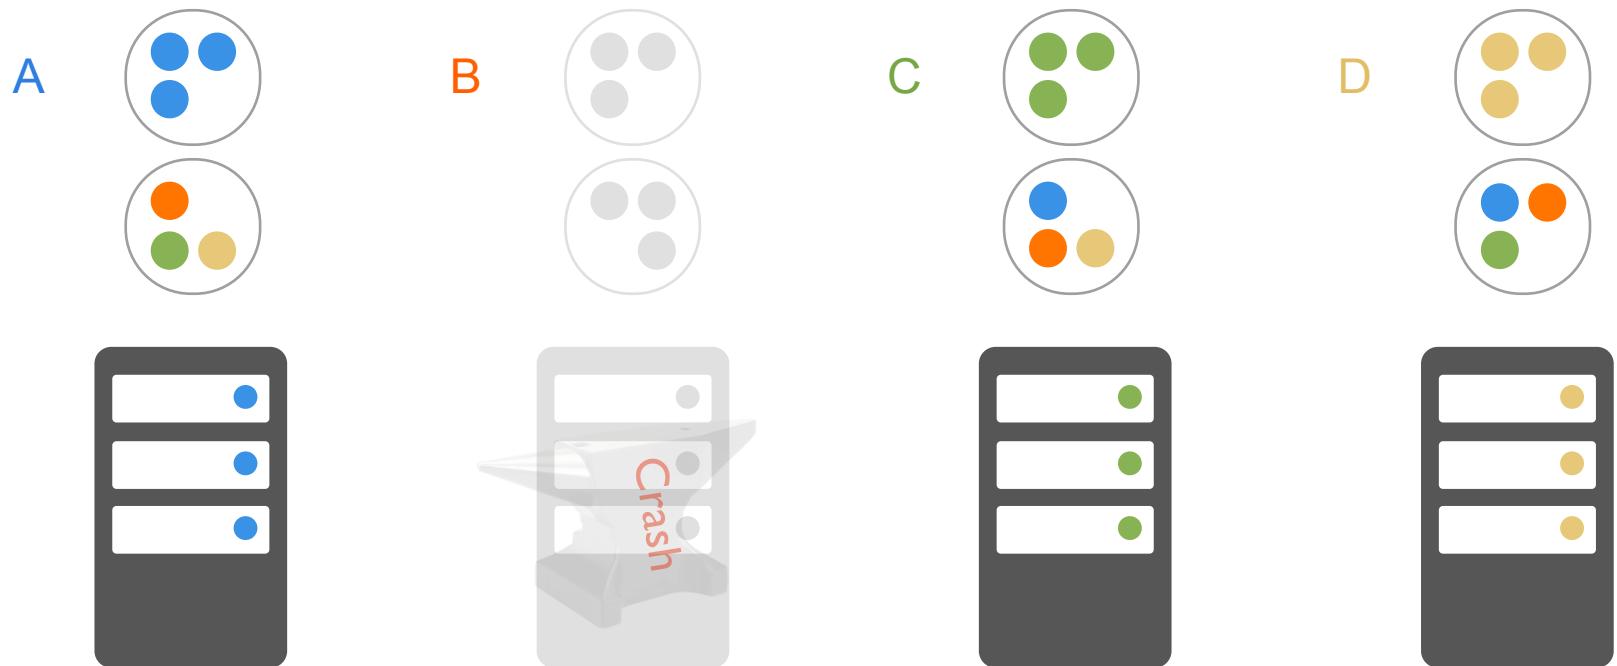


D



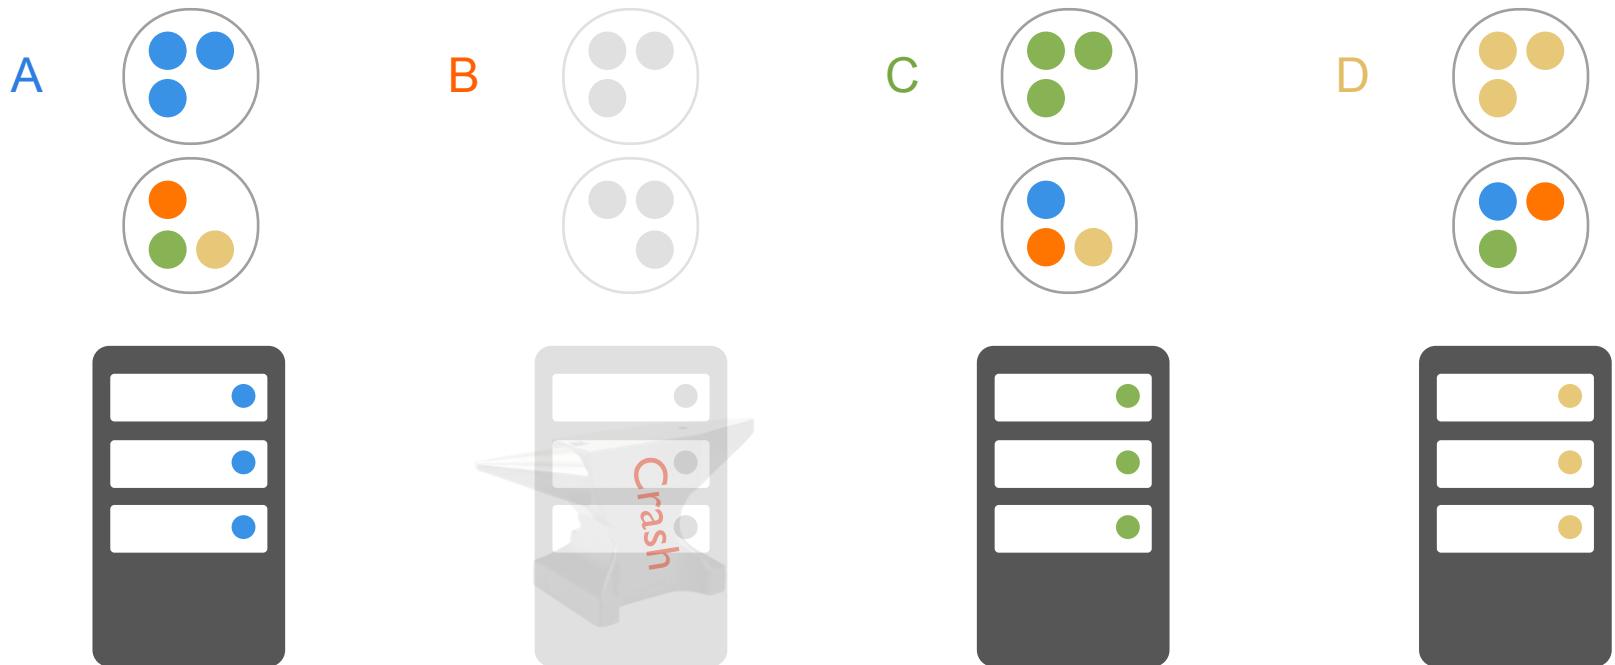


# Backups Are Restored



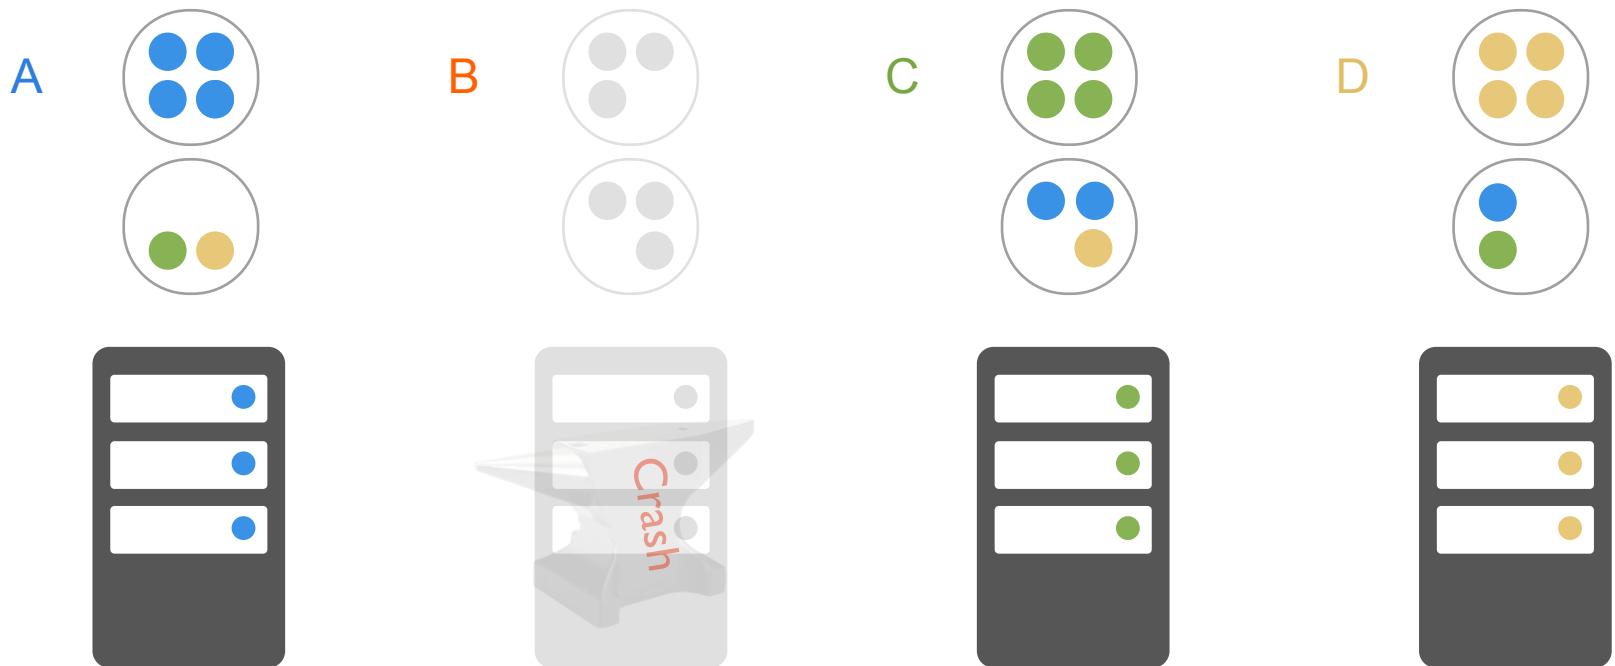


# Backups Are Restored



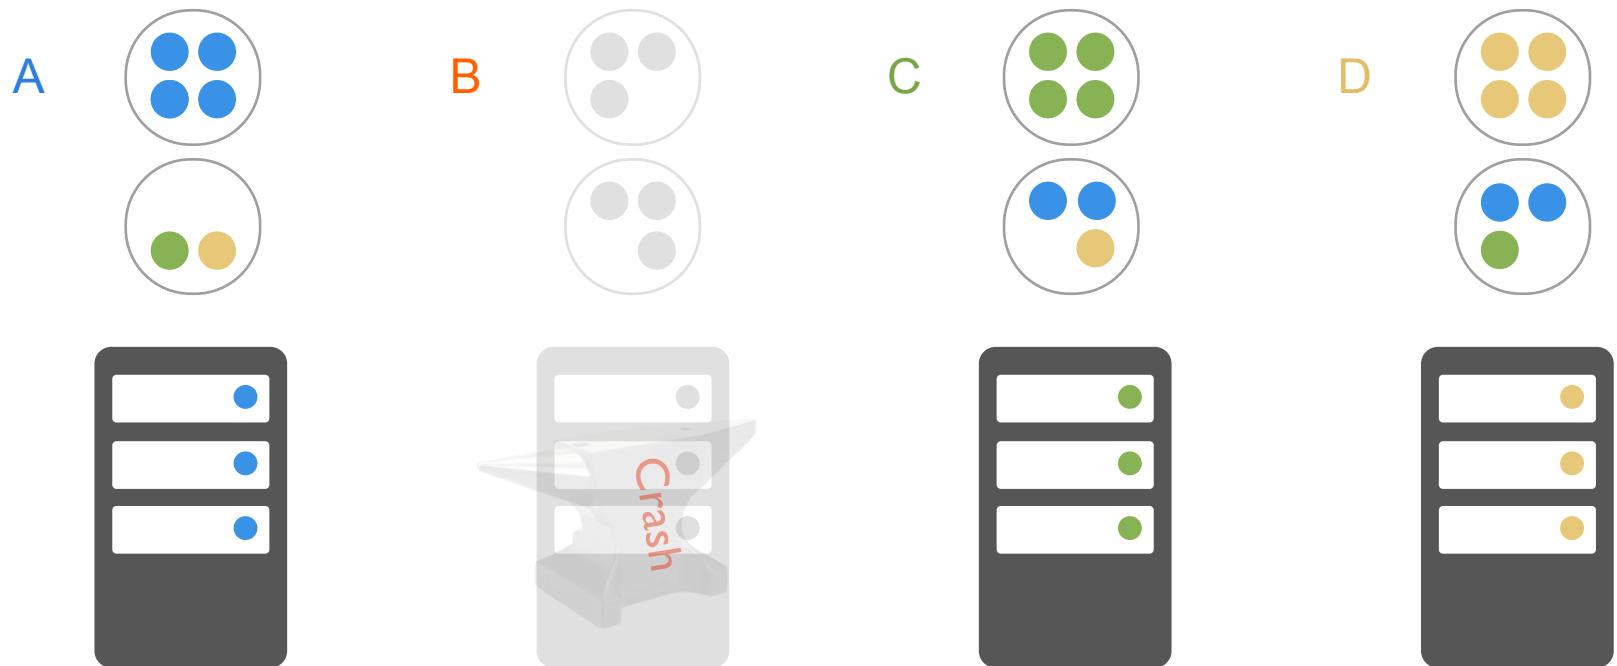


# Backups Are Restored



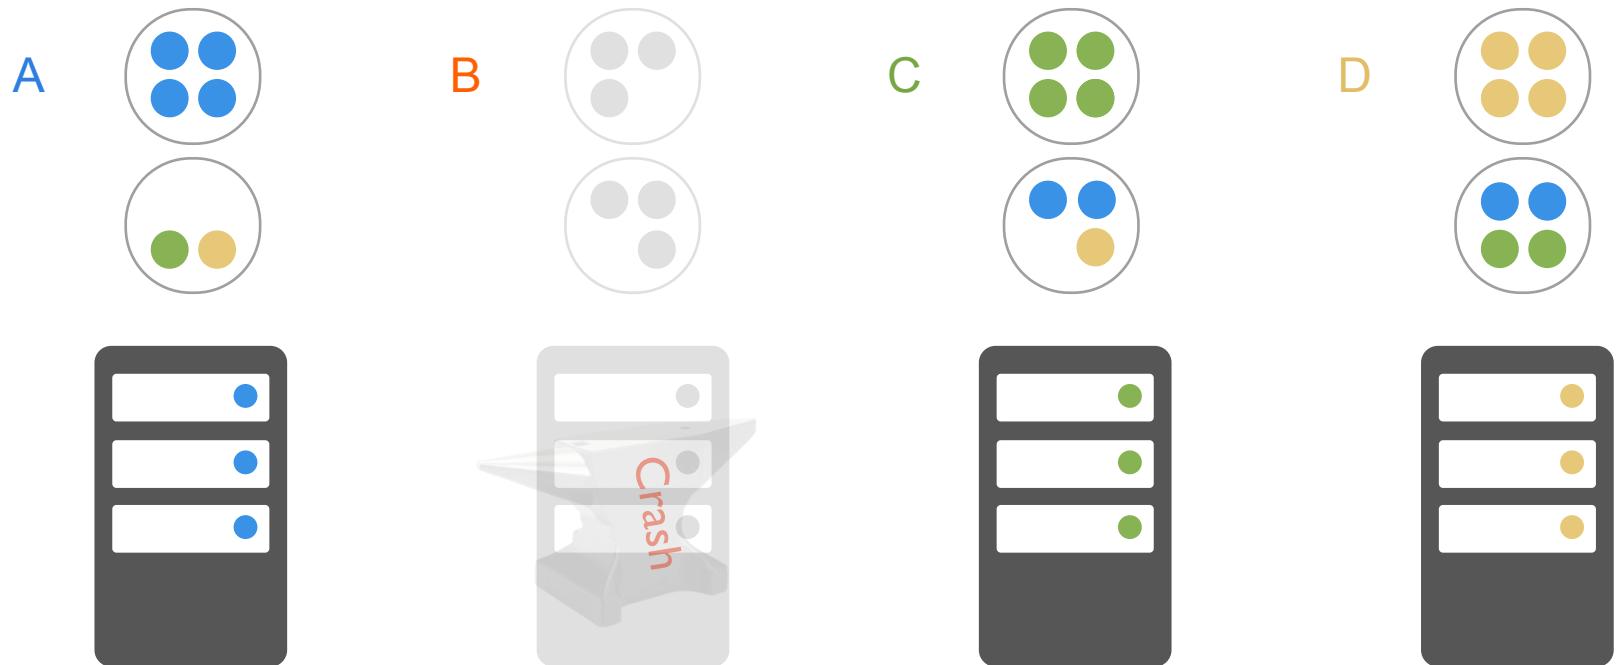


# Backups Are Restored



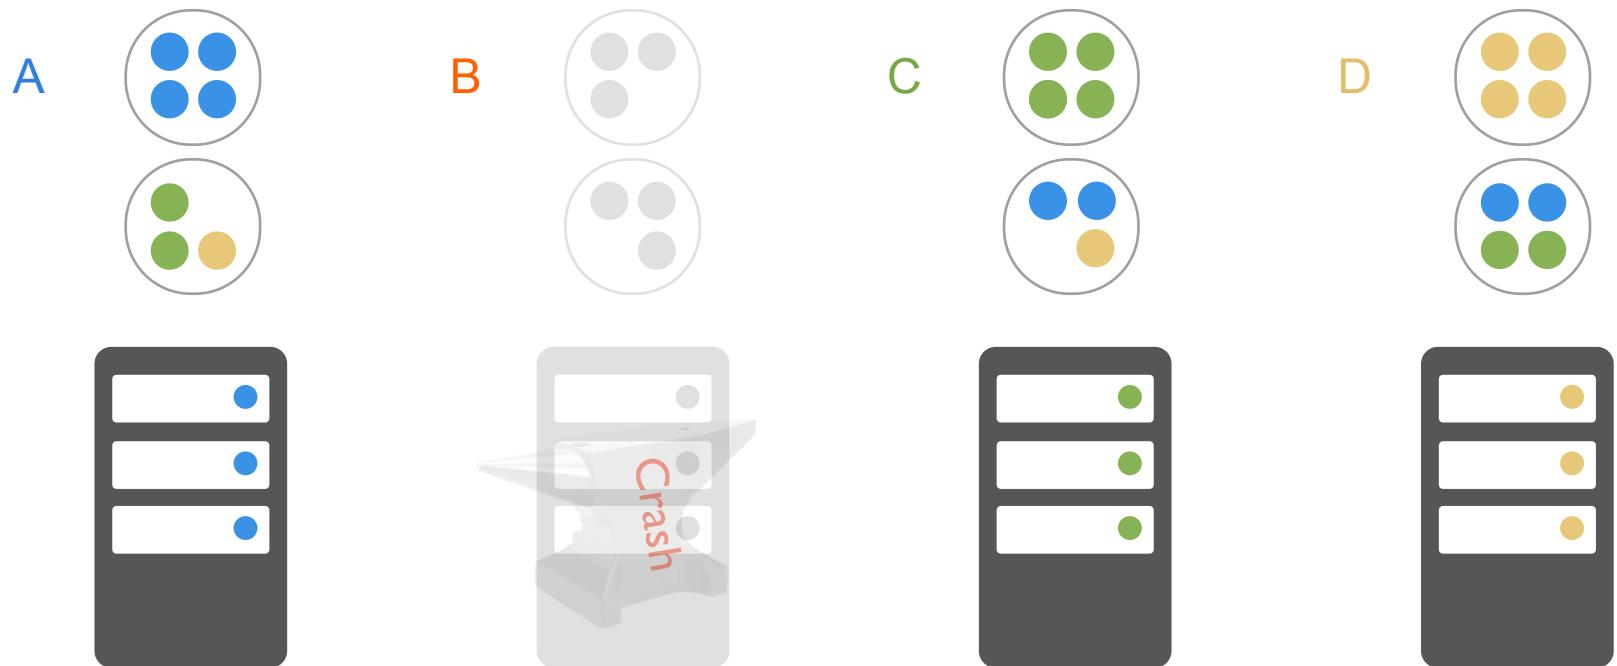


# Backups Are Restored



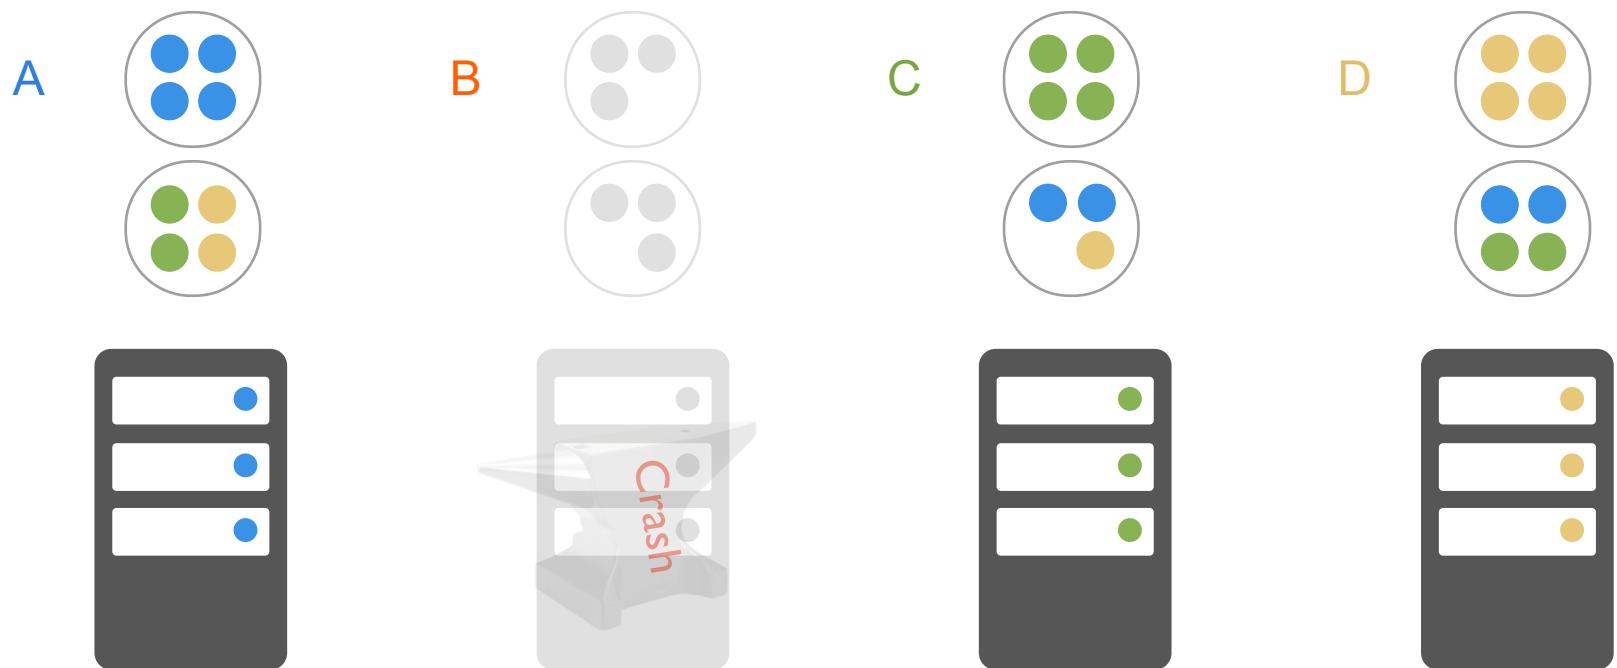


# Backups Are Restored



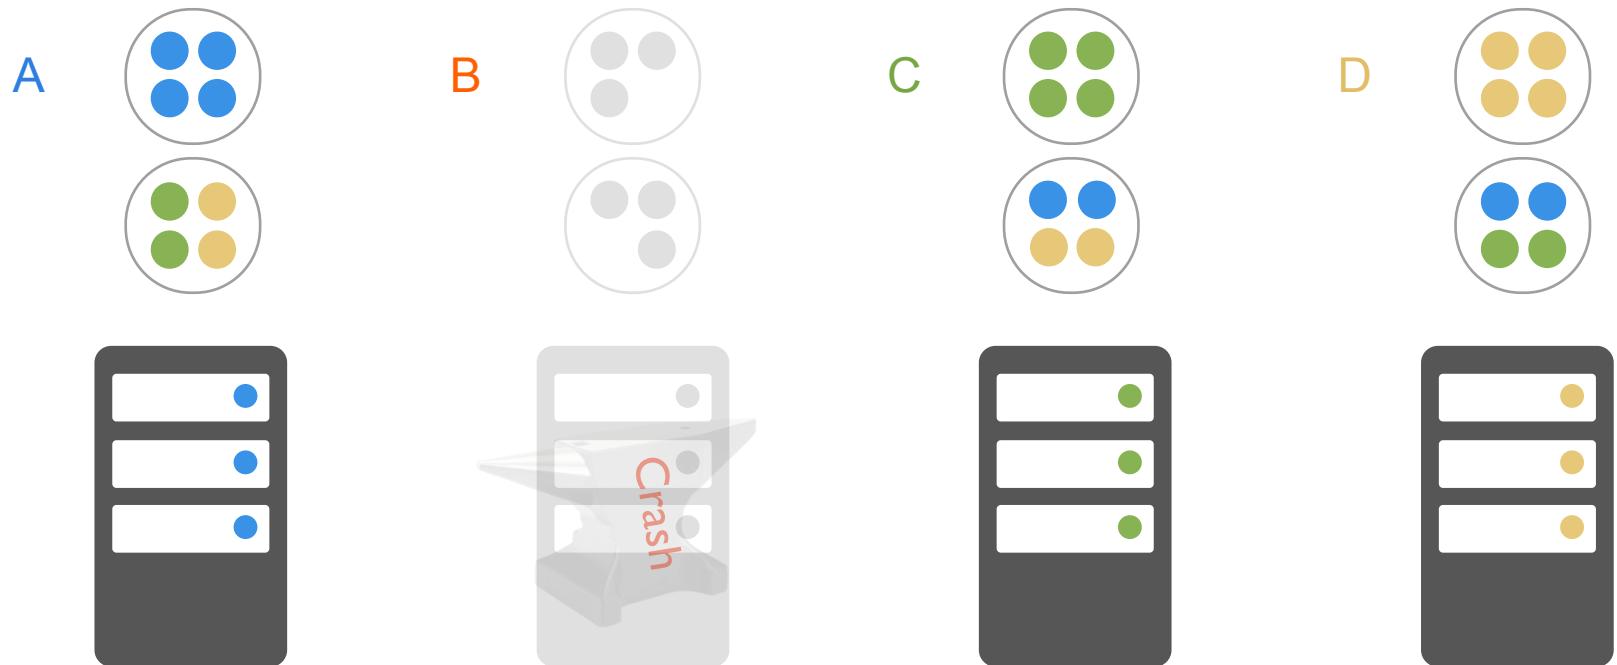


# Backups Are Restored



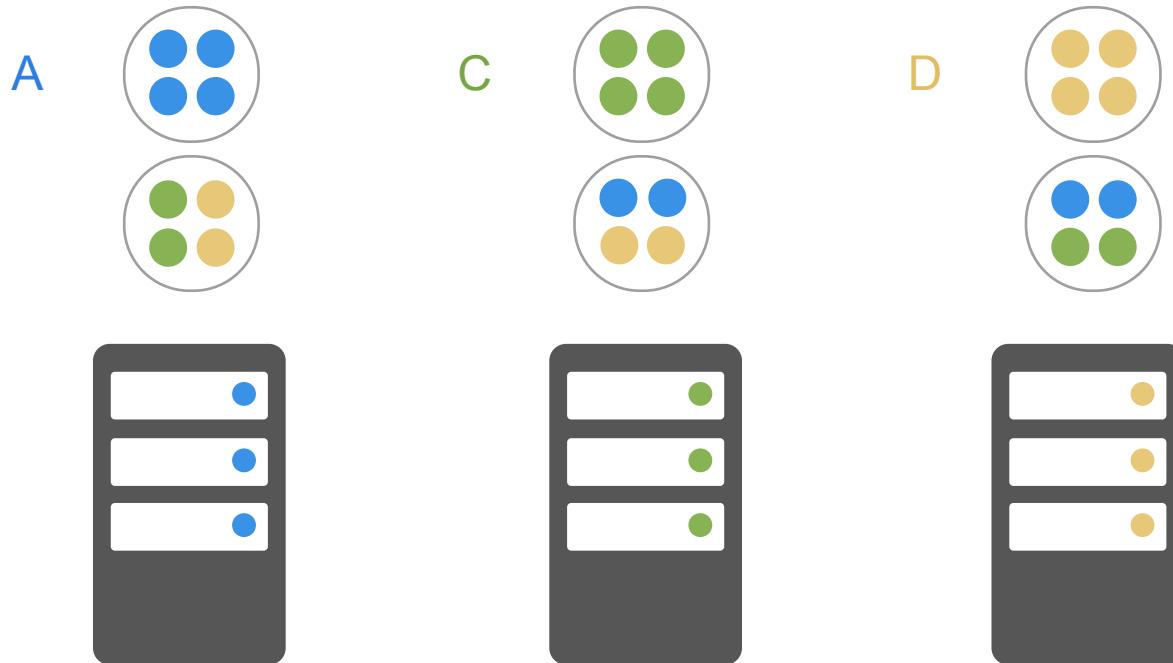


# Backups Are Restored





# Recovery Is Complete



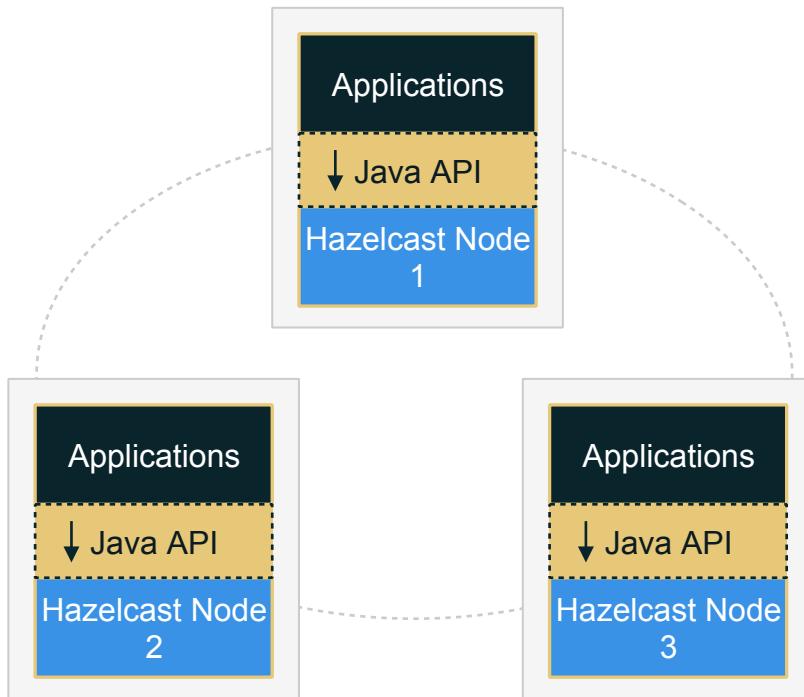


# Deployment Strategies



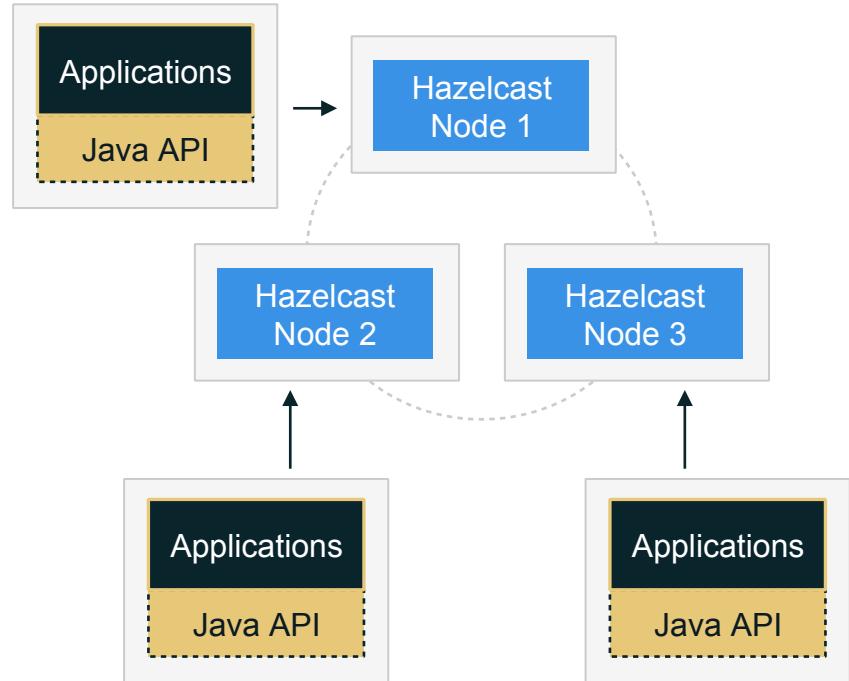
# Deployment Options

## Embedded Hazelcast



Great for early stages of rapid application development and iteration

## Client-Server Mode



Necessary for scale up or scale out deployments – decouples upgrading of clients and cluster for long term TCO



# Easy API

```
// Creating a new Hazelcast node  
HazelcastInstance hz = Hazelcast.newHazelcastInstance();
```

```
// Getting a Map, Queue, Topic, ...  
Map map = hz.getMap("my-map");  
Queue queue = hz.getQueue("my-queue");  
ITopic topic = hz.getTopic("my-topic");
```

```
//Creating a Hazelcast Client  
HazelcastInstance client = Hazelcast.newHazelcastClient();
```

```
// Shutting down the node  
hz.shutdown();
```



# Feature Overview



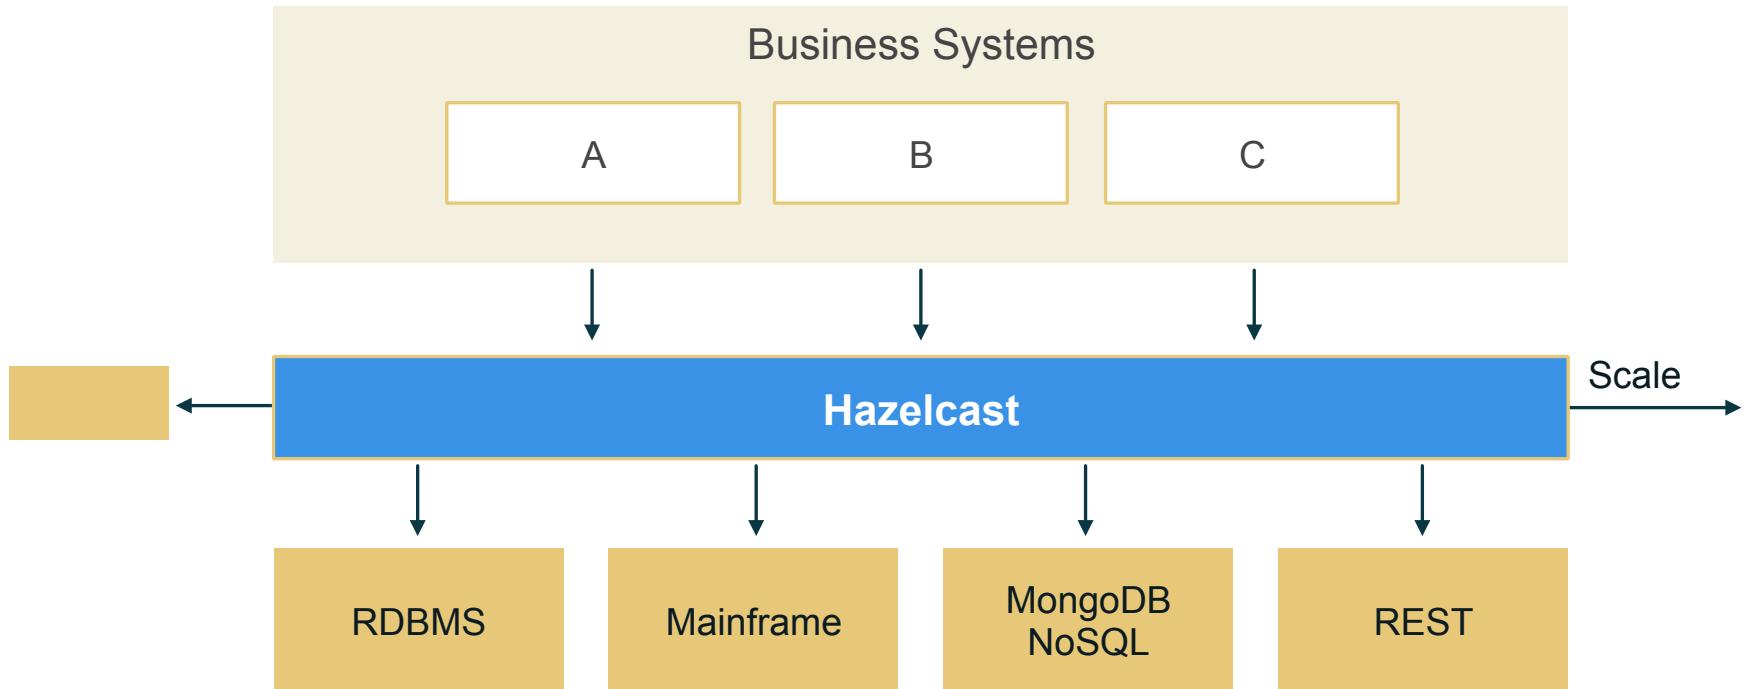
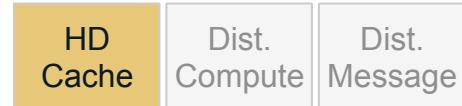
# Easy to Unit Test

```
public class SomeTestCase {  
  
    private HazelcastInstance[] instances;  
  
    @Before  
    public void before() throws Exception {  
        // Multiple instances on the same JVM  
        instances = new HazelcastInstance[2];  
        instances[0] = Hazelcast.newHazelcastInstance();  
        instances[1] = Hazelcast.newHazelcastInstance();  
    }  
  
    @After  
    public void after() throws Exception {  
        Hazelcast.shutdownAll();  
    }  
}
```



# IM Data Store (Caching) Use Case

## Database Caching Use-Case





# IM Data Store (Caching) Features

HD Cache

Dist. Compute

Dist. Message

Java Collection API: Map, List, Set, Queue

JCache

High Density Memory Store

Hibernate 2nd Level Cache

Web Session Replication: Tomcat, Jetty

Predicate API: Indexes, SQL Query

Persistence: Map/Queue Store & Loader. Write Behind/Through

Eviction

Near Cache

Transactions: Local & XA

WAN & DR Replication

Memcached Interface



# Map API

```
interface com.hazelcast.core.IMap<K, V>
    extends java.util.Map, java.util.ConcurrentMap

HazelcastInstance hz = getHazelcastInstance();

//java.util.concurrent.ConcurrentMap implementation
IMap<String, User> hzMap = hz.getMap("users");
hzMap.put("Peter", new User("Peter", "Veentjer"));

hzMap.putIfAbsent("Peter", new User("Peter", "Veentjer"));

//Distributed Lock
hzMap.lock("Peter");

User peter = map.get("Peter");
```



# Persistence API

```
public class MapStorage
    implements MapStore<String, User>, MapLoader<String, User> {
    // Some methods missing ...
    @Override public User load(String key) { return loadValueDB(key); }
    @Override public Set<String> loadAllKeys() { return loadKeysDB(); }
    @Override public void delete(String key) { deleteDB(key); }
    @Override public void store(String key, User value) {
        storeToDatabase(key, value);
    }
}

<map name="users">
    <map-store enabled="true">
        <class-name>com.hazelcast.example.MapStorage</class-name>
        <write-delay-seconds>0</write-delay-seconds>
    </map-store>
</map>
```



# JCache API

```
// Retrieve the CachingProvider which is automatically baced by  
// the chosen Hazelcast server or client provider  
CachingProvider cachingProvider = Caching.getCachingProvider();  
  
// Create a CacheManager  
CacheManager cacheManager = cachingProvider.getCacheManager();  
  
// Cache<String, String> cache = cacheManager  
//     .getCache( name, String.class, String.class );  
  
// Create a simple but typesafe configuration for the cache  
  
CompleteConfiguration<String, String> config =  
    new MutableConfiguration<String, String>()  
        .setTypes( String.class, String.class );
```



# JCache API

```
// Create and get the cache
Cache<String, String> cache = cacheManager
    .createCache( "example", config );
// Alternatively to request an already existing cache
Cache<String, String> cache = cacheManager
    .getCache( name, String.class, String.class );

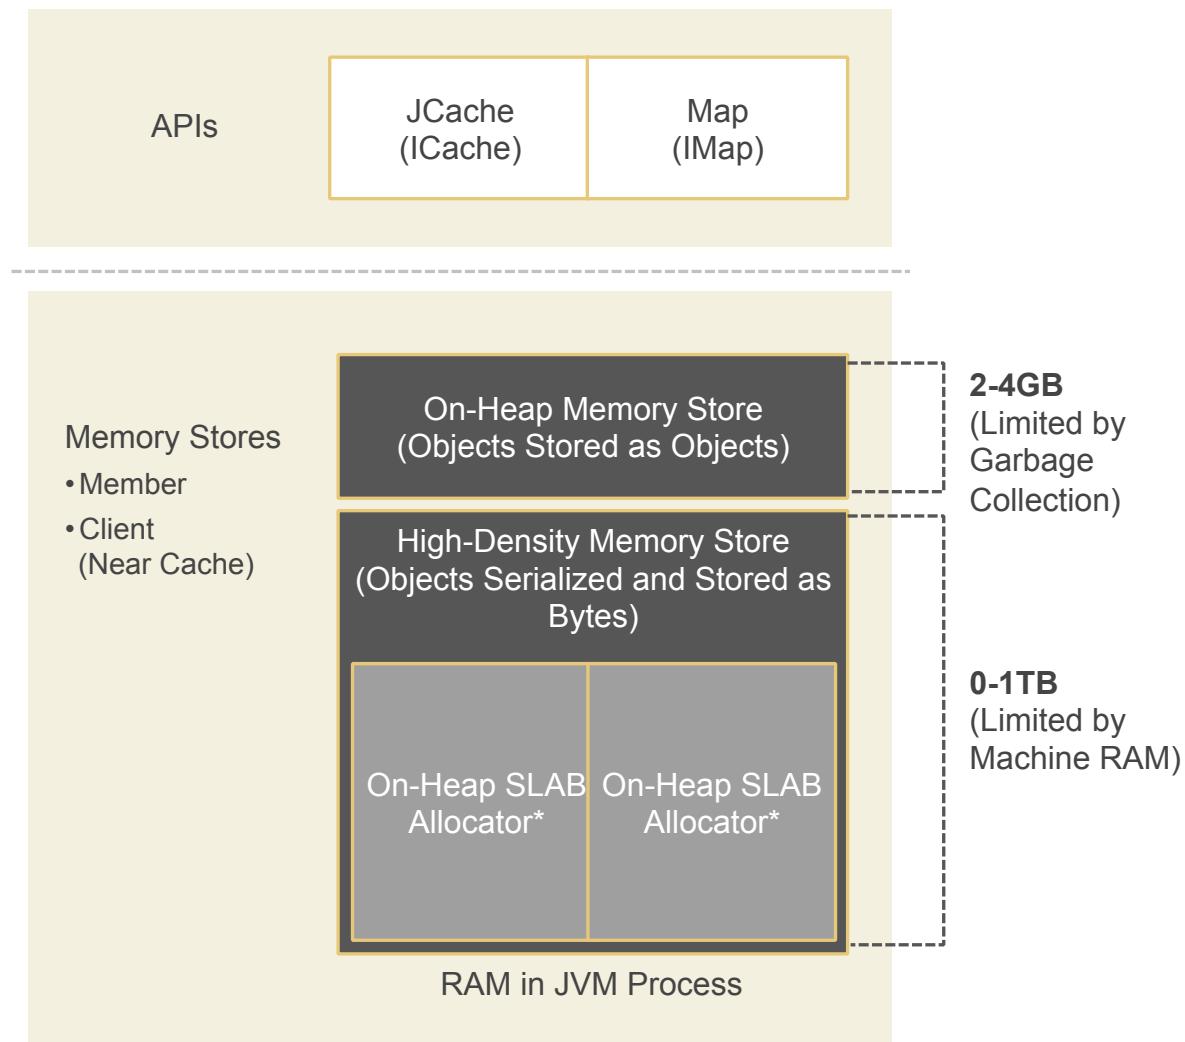
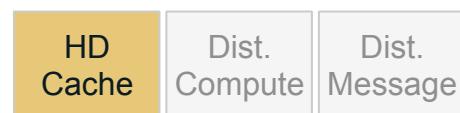
// Put a value into the cache
cache.put( "world", "Hello World" );

// Retrieve the value again from the cache
String value = cache.get( "world" );

System.out.println( value );
```



# High Density Caching



\* coming in 3.6



# On Heap Vs. High-Density Memory Management

HD Cache    Dist. Compute    Dist. Message

## On Heap Memory

**0 MB**

**3.9 GB**

**9 (4900 ms)**

**31 (4200 ms)**

**Native**

**Heap Storage**

**Major GC**

**Minor GC**

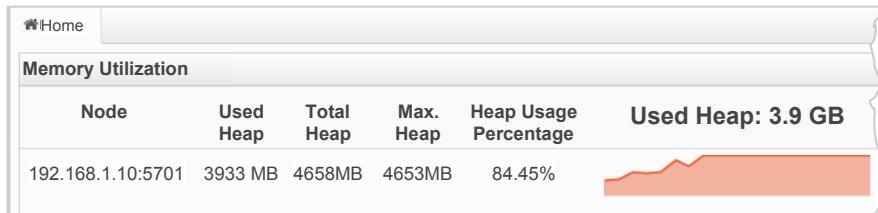
## HD Memory v2

**3.3 GB**

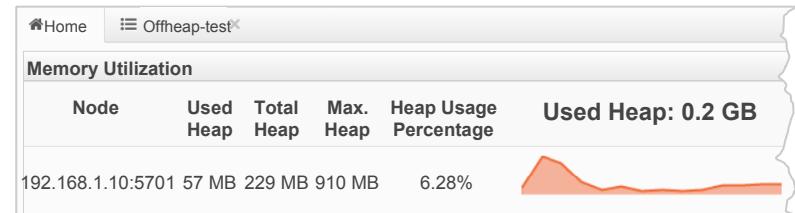
**0.6 GB**

**0 (0 ms)**

**356 (349 ms)**



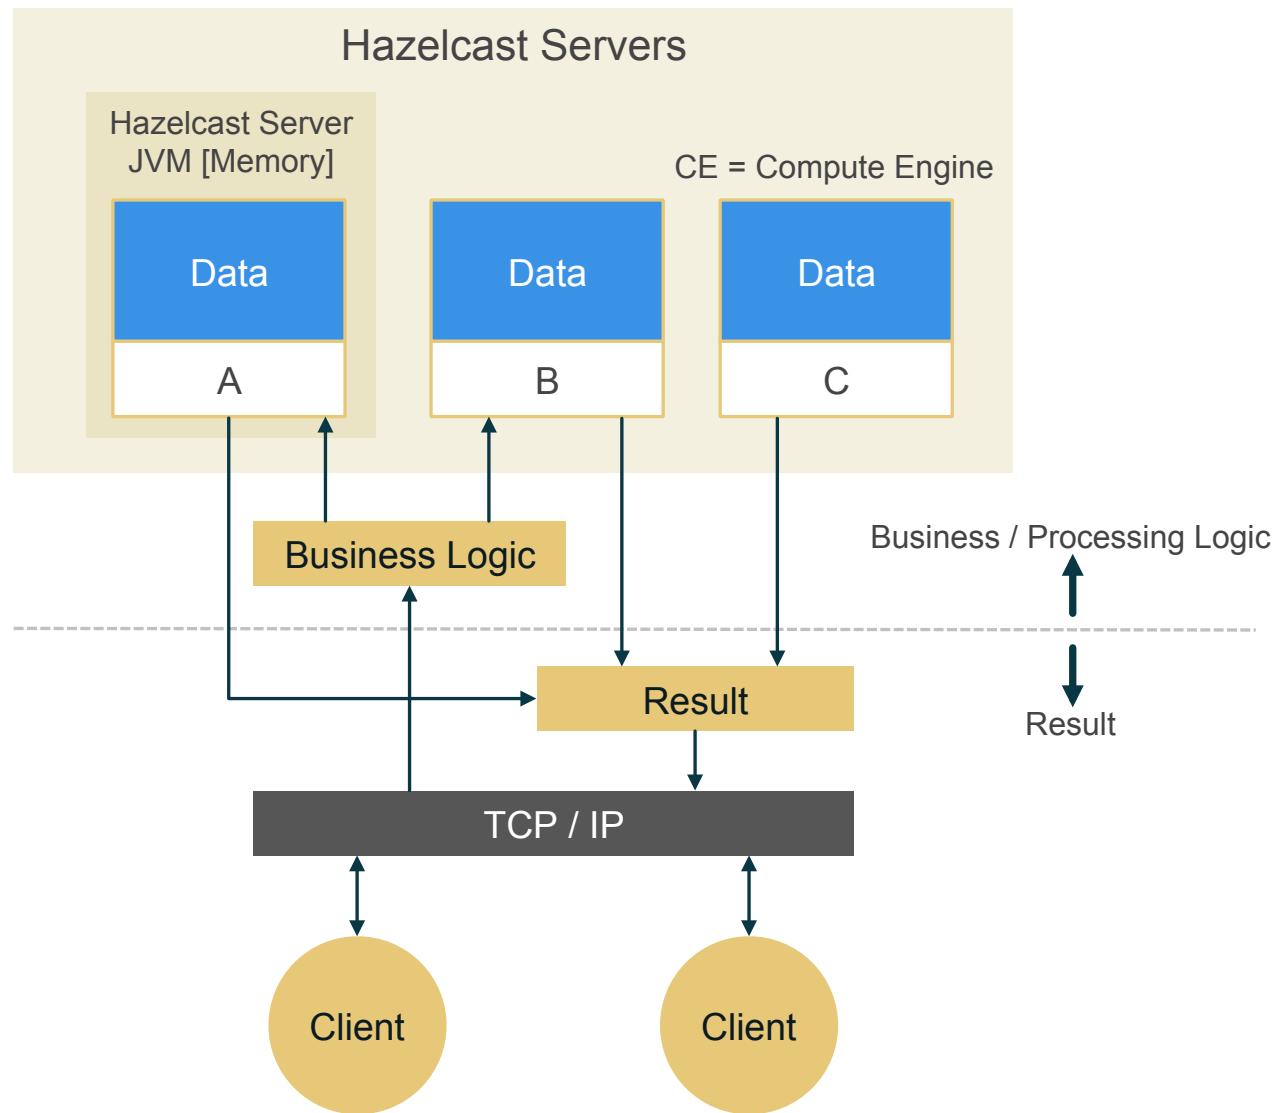
Example: On Heap Memory



Example: HD Memory v2



# IM Distributed Computing Use Case





# IM Distributed Computing Feature



Java Concurrency API  
(Lock, Semaphore, AtomicLong, AtomicReference, Executor Service, Blocking Queue)

Entry and Item Listeners

Entry Processor

Aggregators

Map/Reduce

Data Affinity

Continues Query

Map Interceptors

Delta Update



# Lock API

```
HazelcastInstance hz = getHazelcastInstance();
```

```
// Distributed Reentrant
```

```
Lock lock = hz.getLock("myLock");
lock.lock();
try {
    // Do something
} finally {
    lock.unlock();
}
```



# Executor Service API

```
public interface com.hazelcast.core.IExecutorService  
    extends java.util.concurrent.ExecutorService  
  
HazelcastInstance hz = getHazelcastInstance();  
  
//java.util.concurrent.ExecutorService implementation  
IExecutorService es = hz.getExecutorService("name");  
es.executeOnAllMembers(buildRunnable());  
es.executeOnKeyOwner(buildRunnable(), "Peter");  
es.execute(buildRunnable());  
  
Map<..> futures = es.submitToAllMembers(buildCallable());  
Future<..> future = es.submitToKeyOwner(buildCallable(), "Peter");  
  
es.submitToAllMembers(buildCallable(), buildCallback());  
es.submitToKeyOwner(buildCallable(), "Peter", buildCallback());
```



# Map/Reduce API

```
HazelcastInstance hz = getHazelcastInstance();

Map users = hz.getMap("users");
JobTracker tracker = hz.getJobTracker("default");

KeyValueSource source = KeyValueSource.fromMap(users);
Job job = tracker.newJob(source);

ICompleteFuture future = job.mapper(new MyMapper())
    .reducer(new MyReducer())
    .submit();

Map result = future.get();
```



# Aggregations API

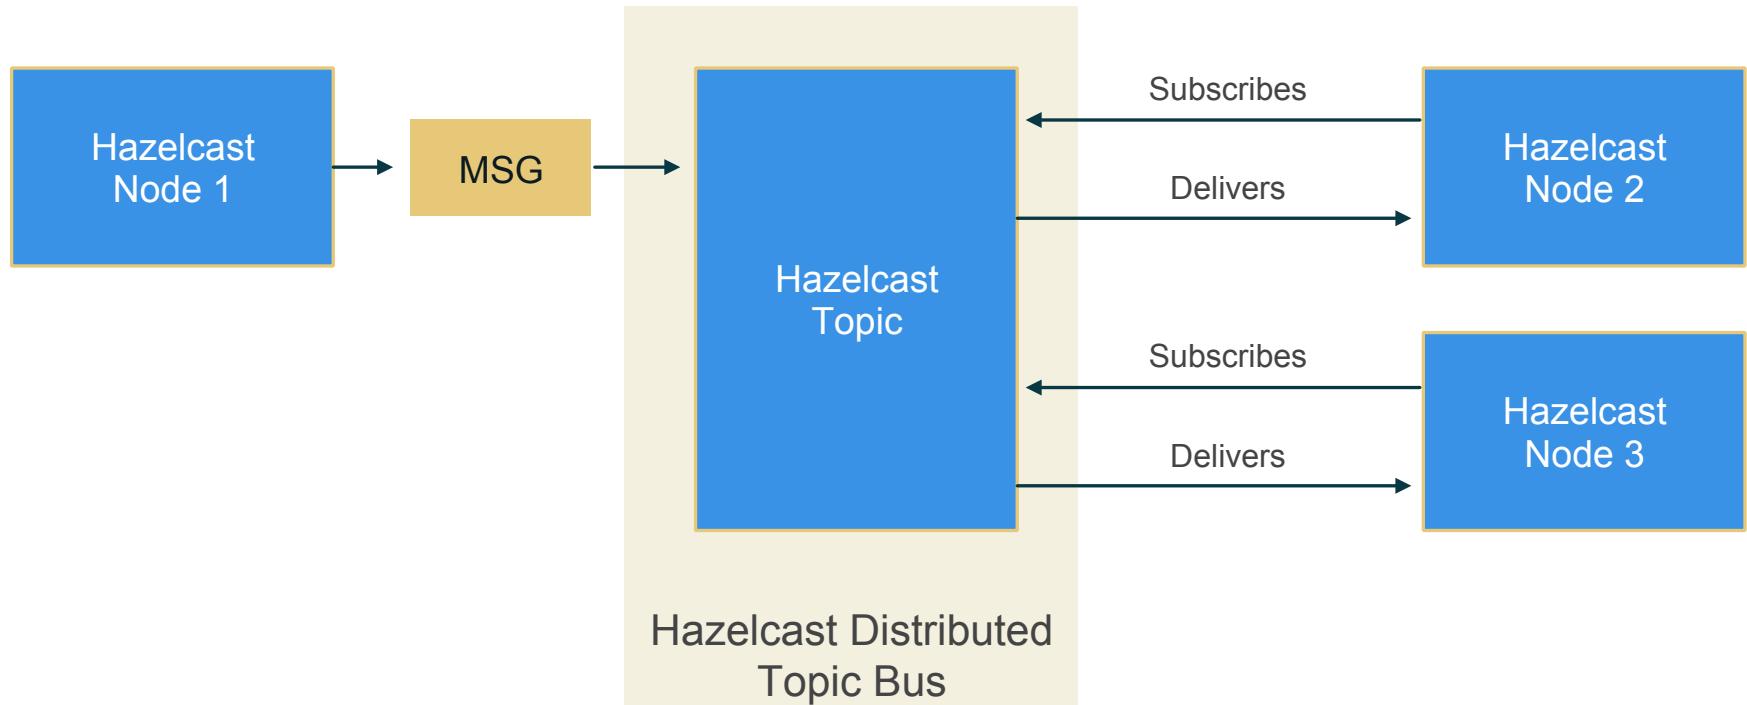
```
HazelcastInstance hz = getHazelcastInstance();

Map users = hz.getMap("users");

int sum = users.aggregate(
    Supplier.all((user) -> user.getSalary()),
    Aggregations.longSum()
);
```

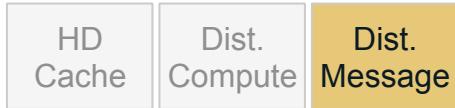


# IM Distributed Messaging Use Case





# IM Distributed Messaging Features



Queue

Topic (Pub/Sub)

Event Listeners

Ring Buffers



# Queue API

```
interface com.hazelcast.core.IQueue<E>
    extends java.util.concurrent.BlockingQueue

HazelcastInstance hz = getHazelcastInstance();

//java.util.concurrent.BlockingQueue implementation
IQueue<Task> queue = hz.getQueue("tasks");

queue.offer(newTask());
queue.offer(newTask(), 500, TimeUnit.MILLISECONDS);

Task task = queue.poll();
Task task = queue.poll(100, TimeUnit.MILLISECONDS);
Task task = queue.take();
```



# Topic API

```
public class Example implements MessageListener<String> {
    public void sendMessage {
        HazelcastInstance hz = getHazelcastInstance();
        ITopic<String> topic = hz.getTopic("topic");
        topic.addMessageListener(this);
        topic.publish("Hello World");
    }

    @Override
    public void onMessage(Message<String> message) {
        System.out.println("Got message: " + message.getMessageObject());
    }
}
```



# Thank you

@fuadm, @hazelcast

[hazelcast@googlegroups.com](mailto:hazelcast@googlegroups.com)

<http://www.hazelcast.com>

<http://github.com/hazelcast/hazelcast>