



Thinking Distributed: The Hazelcast Way

RAHUL GUPTA

SR SOLUTIONS ARCHITECT

About Me

Rahul Gupta



Senior Solutions Architect

Worked with Terracotta

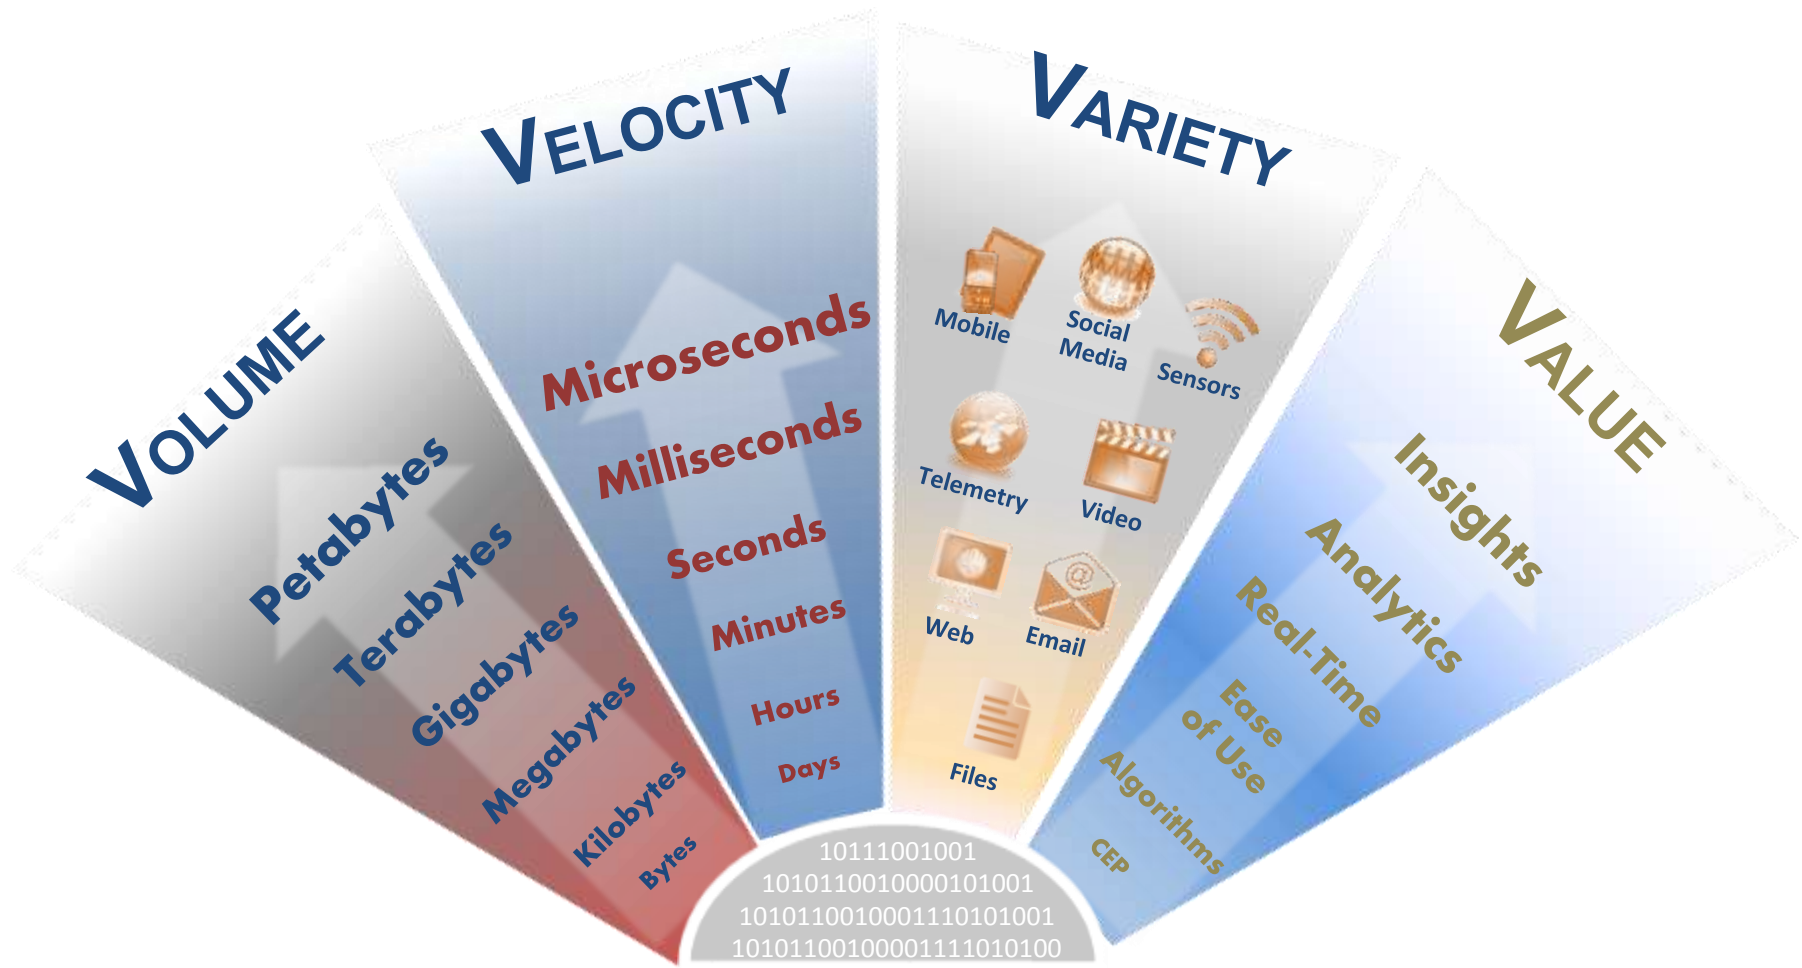
In-memory Distributed Systems since 2009

Java Programmer since 1998

rahul@hazelcast.com

follow me [@wildnez](#)

Challenges



Why Hazelcast?



- **Scale-out Computing** enables cluster capacity to be increased or decreased on-demand



- **Resilience** with automatic recovery from member failures without losing data while minimizing performance impact on running applications

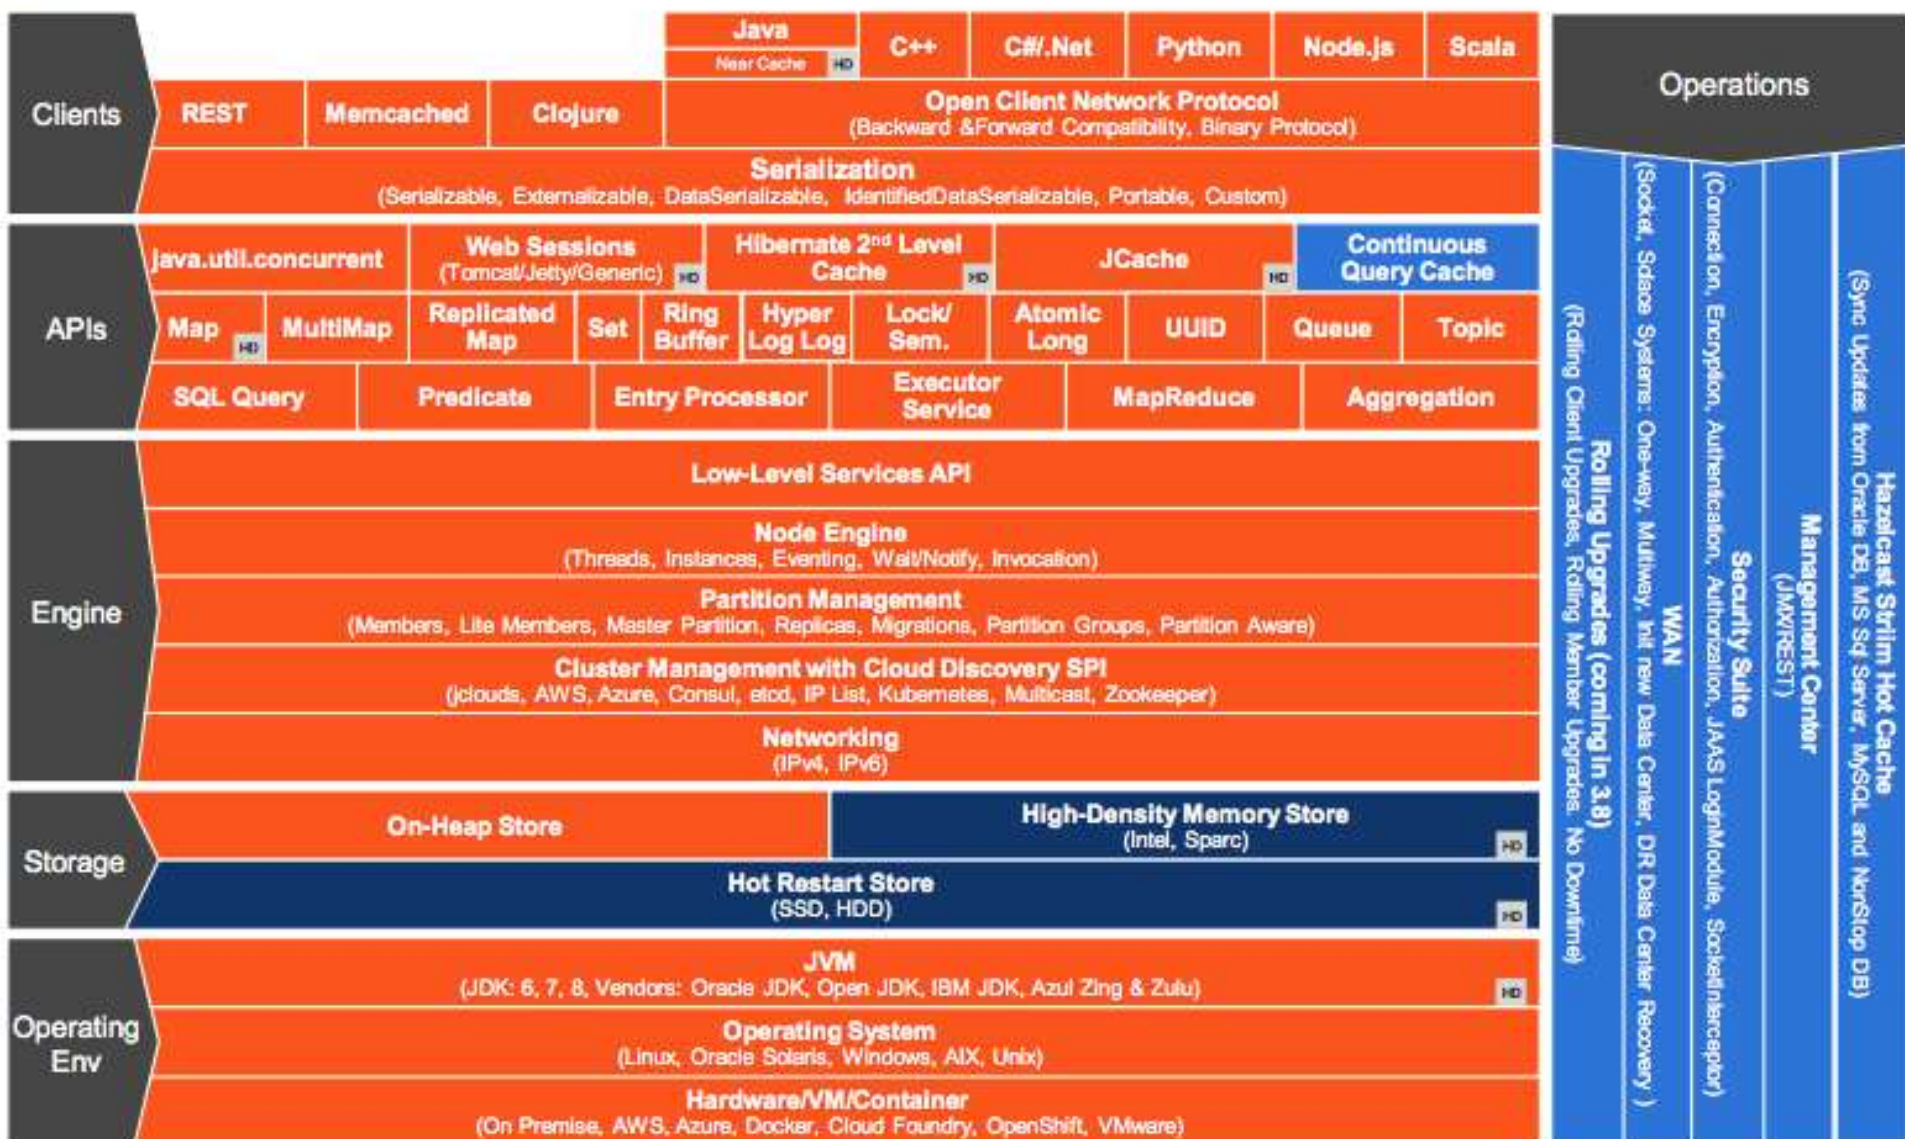


- **Programming Model** provides a way for developers to easily program a cluster application as if it is a single process



- **Fast Application Performance** enables very large data sets to be held in main memory for real-time performance

Hazelcast IMDG 3.7/3.8



 Open Source
  Enterprise Edition
  Enterprise HD Edition
  HD-Enabled Feature




Ecosystem Traction

Dozens of Commercial and Open Source Projects Embed Hazelcast




In Memory Data Grid



In Memory
Data **Storage**

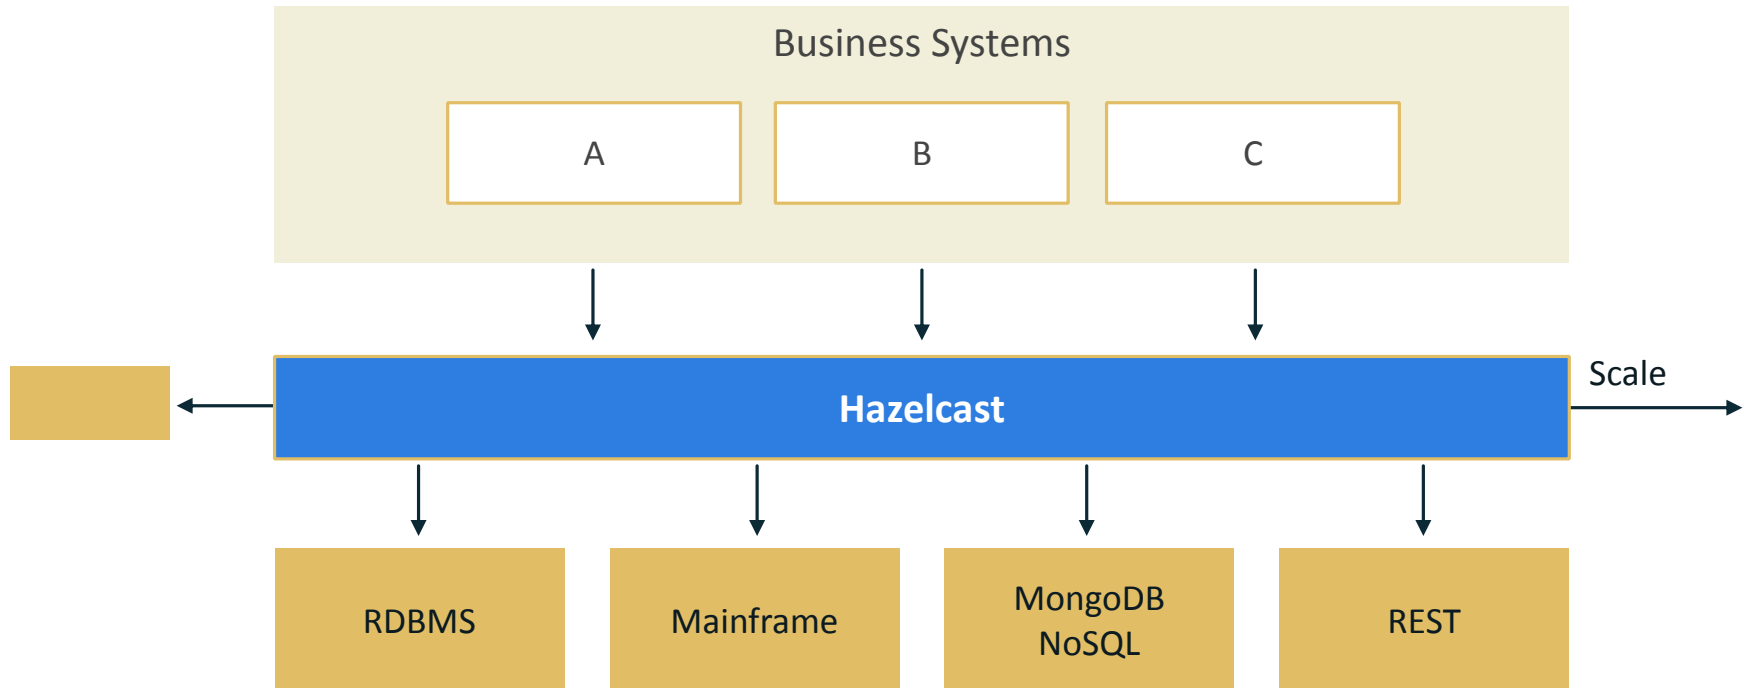


In Memory
Data **Computing**



In Memory
Data **Messaging**

In-Memory Caching



java.util.concurrent.ConcurrentMap

```
import java.util.concurrent.ConcurrentHashMap;  
import java.util.concurrent.ConcurrentMap;  
  
public static void main(String[] args) {  
    ConcurrentMap<Integer, String> map = new ConcurrentHashMap<>();  
    map.put(1, "Paris");  
    map.put(2, "London");  
    map.put(3, "San Francisco");  
  
    String oldValue = map.remove(2);  
}
```



Distributed Map

```
import java.util.concurrent.ConcurrentMap;
import com.hazelcast.core.Hazelcast;
import com.hazelcast.core.HazelcastInstance;

public static void main(String[] args) {
    HazelcastInstance h = Hazelcast.newHazelcastInstance();

    ConcurrentMap<Integer, String> map = h.getMap("myMap");
    map.put(1, "Paris");
    map.put(2, "London");
    map.put(3, "San Francisco");

    String oldValue = map.remove(2);
}
```



DEMO



Persistence API

```
public class MapStorage
    implements MapStore<String, User>, MapLoader<String, User> {
    // Some methods missing ...
    @Override public User load(String key) { return loadValueDB(key); }
    @Override public Set<String> loadAllKeys() { return loadKeysDB(); }
    @Override public void delete(String key) { deleteDB(key); }
    @Override public void store(String key, User value) {
        storeToDatabase(key, value);
    }
}
```

```
<map name="users">
  <map-store enabled="true">
    <class-name>com.hazelcast.example.MapStorage</class-name>
    <write-delay-seconds>0</write-delay-seconds>
  </map-store>
</map>
```



JCache API

```
// Retrieve the CachingProvider which is automatically baced by
// the chosen Hazelcast server or client provider
CachingProvider cachingProvider = Caching.getCachingProvider();

// Create a CacheManager
CacheManager cacheManager = cachingProvider.getCacheManager();

// Cache<String, String> cache = cacheManager
//   .getCache( name, String.class, String.class );

// Create a simple but typesafe configuration for the cache

CompleteConfiguration<String, String> config =
    new MutableConfiguration<String, String>()
        .setTypes( String.class, String.class );
```



JCache API

```
// Create and get the cache
```

```
Cache<String, String> cache = cacheManager  
    .createCache( "example", config );
```

```
// Alternatively to request an already existing cache
```

```
Cache<String, String> cache = cacheManager  
    .getCache( name, String.class, String.class );
```

```
// Put a value into the cache
```

```
cache.put( "world", "Hello World" );
```

```
// Retrieve the value again from the cache
```

```
String value = cache.get( "world" );
```

```
System.out.println( value );
```



Eviction

- Unless deleted, entries remain in the map.
- Use eviction policies to prevent OOM situations.
- Eviction Triggers run LFU or LRU.

time-to-live

max-idle-seconds

max-size (PER_NODE(number of entries),
PER_PARTITION(number of entries),
USED_HEAP_SIZE(mb),
USED_HEAP_PERCENTAGE(mb))

- Setting eviction-percentage removes that % of entries when eviction is triggered.



More Distributed Structures

Features	Description
MultiMap	Store multiple values against one Key
Replicated Map	Cluster wide replication, all entries on all nodes. Good for read heavy use cases
Near Cache	Map Entries on Client/Application local memory
RingBuffer	Stores data in a ring-like structure, like a circular array with given capacity



IM Data Store (Caching) Features

Java Collection API: Map, List, Set, Queue
Jcache
High Density Memory Store
Hibernate 2nd Level Cache
Web Session Replication: Tomcat, Jetty
Predicate API: Indexes, SQL Query
Persistence: Map/Queue Store & Loader. Write Behind/Through
Spring Compliance
Transactions: Local & XA
WAN & DR Replication



Java: Will it make the cut?

Garbage Collection limits heap usage. G1 and Balanced aim for <100ms at 10GB.

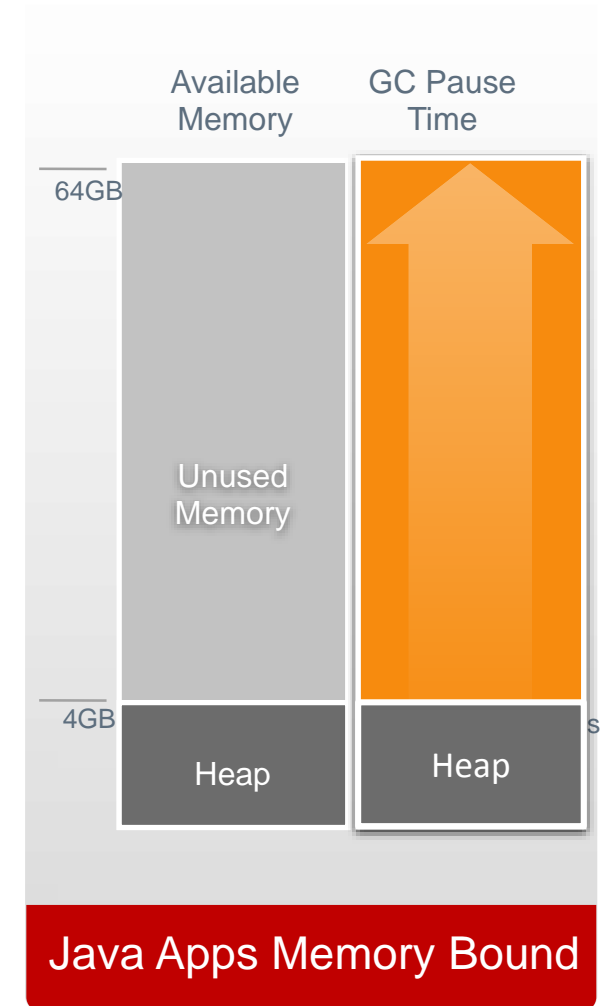


Off-Heap Storage

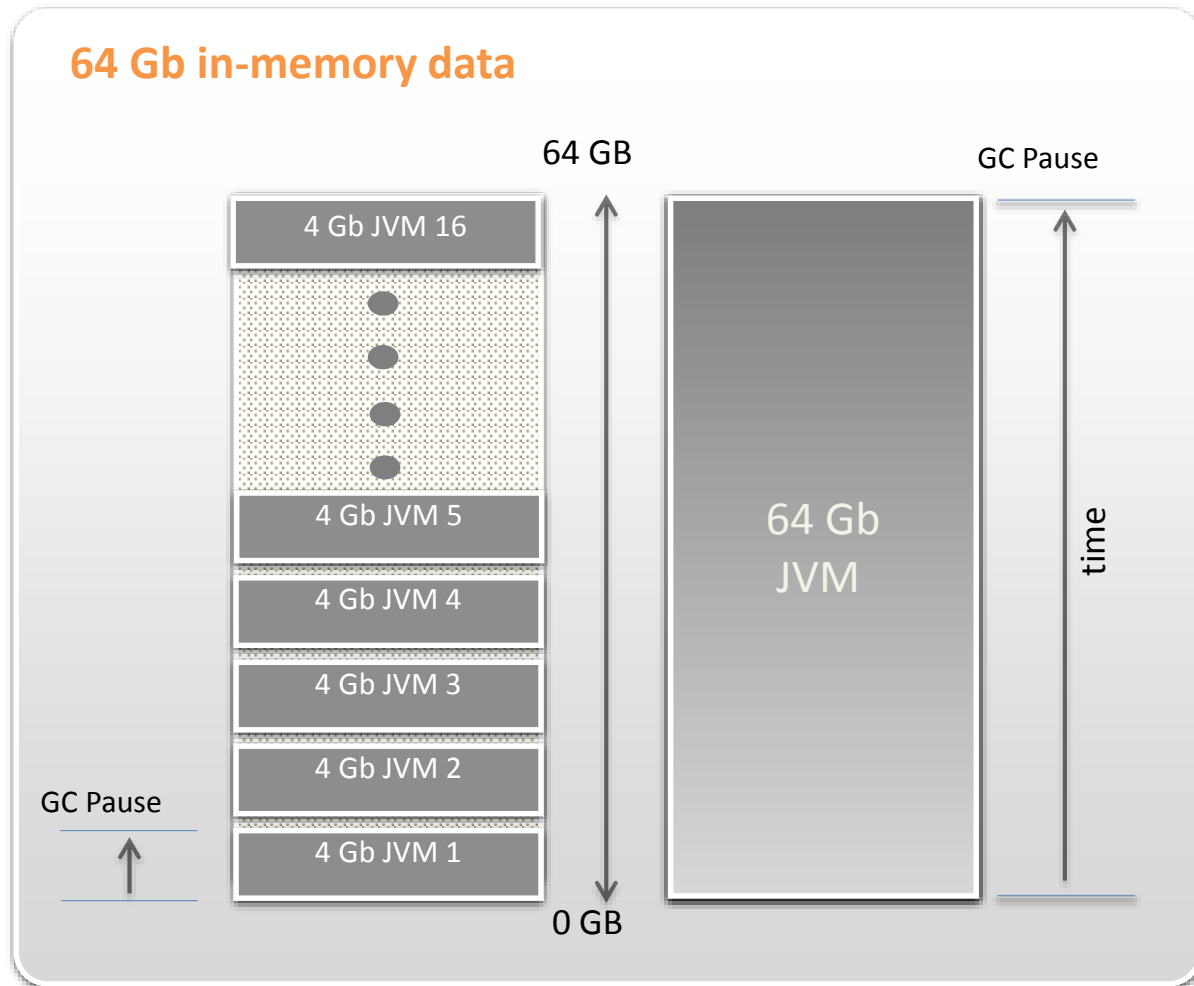
No low-level CPU access



Java is challenged as an infrastructure language despite its newly popular usage for this




Standard Impediments of Caching

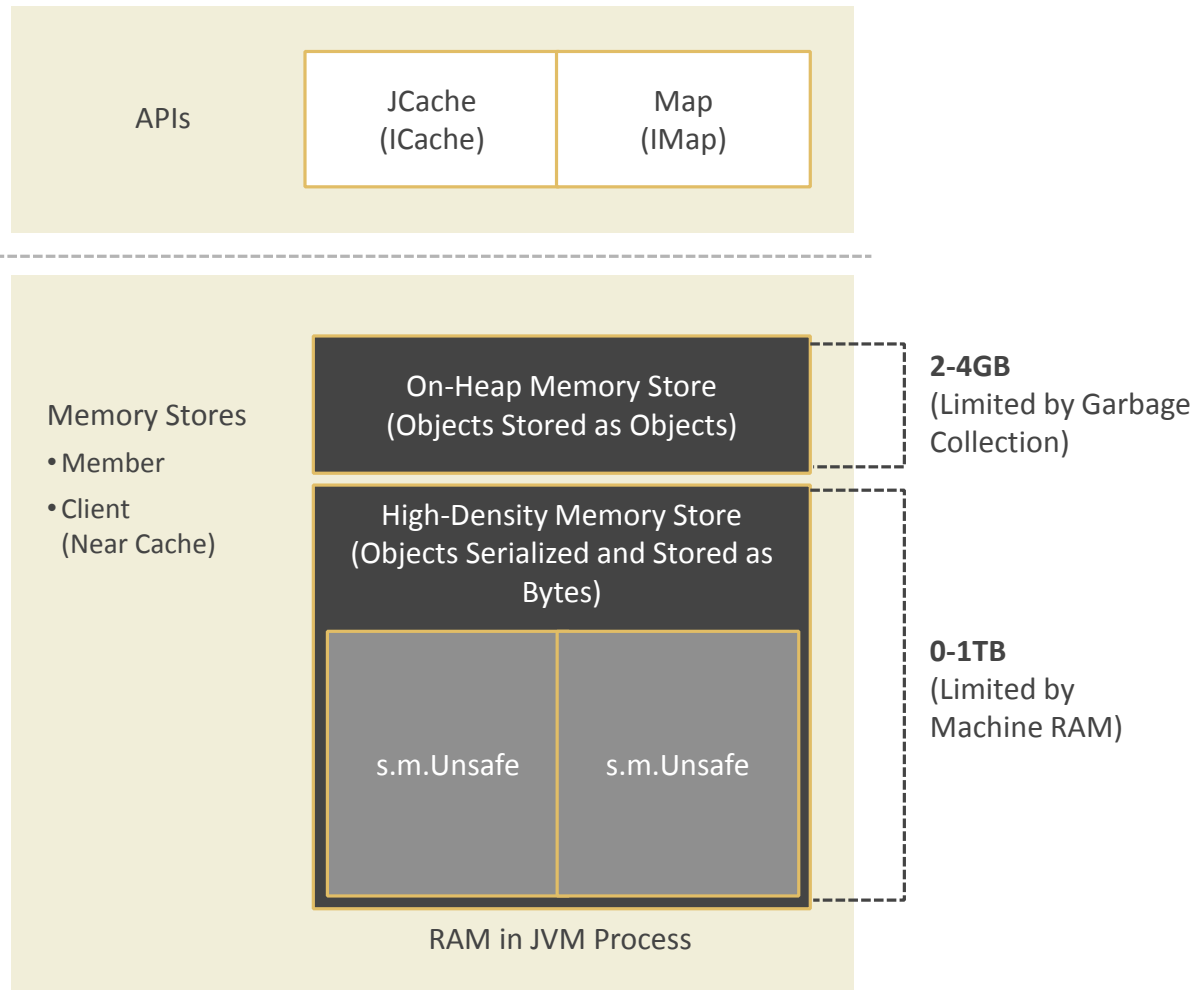


Caching with HD Memory

64 Gb in-memory data



HD Memory



On Heap Vs. High-Density Memory Management

On Heap Memory

0 MB

3.9 GB

9 (4900 ms)

31 (4200 ms)

HD

Heap Storage

Major GC

Minor GC


HD Memory

3.3 GB


0.6 GB

0 (0 ms)

356 (349 ms)

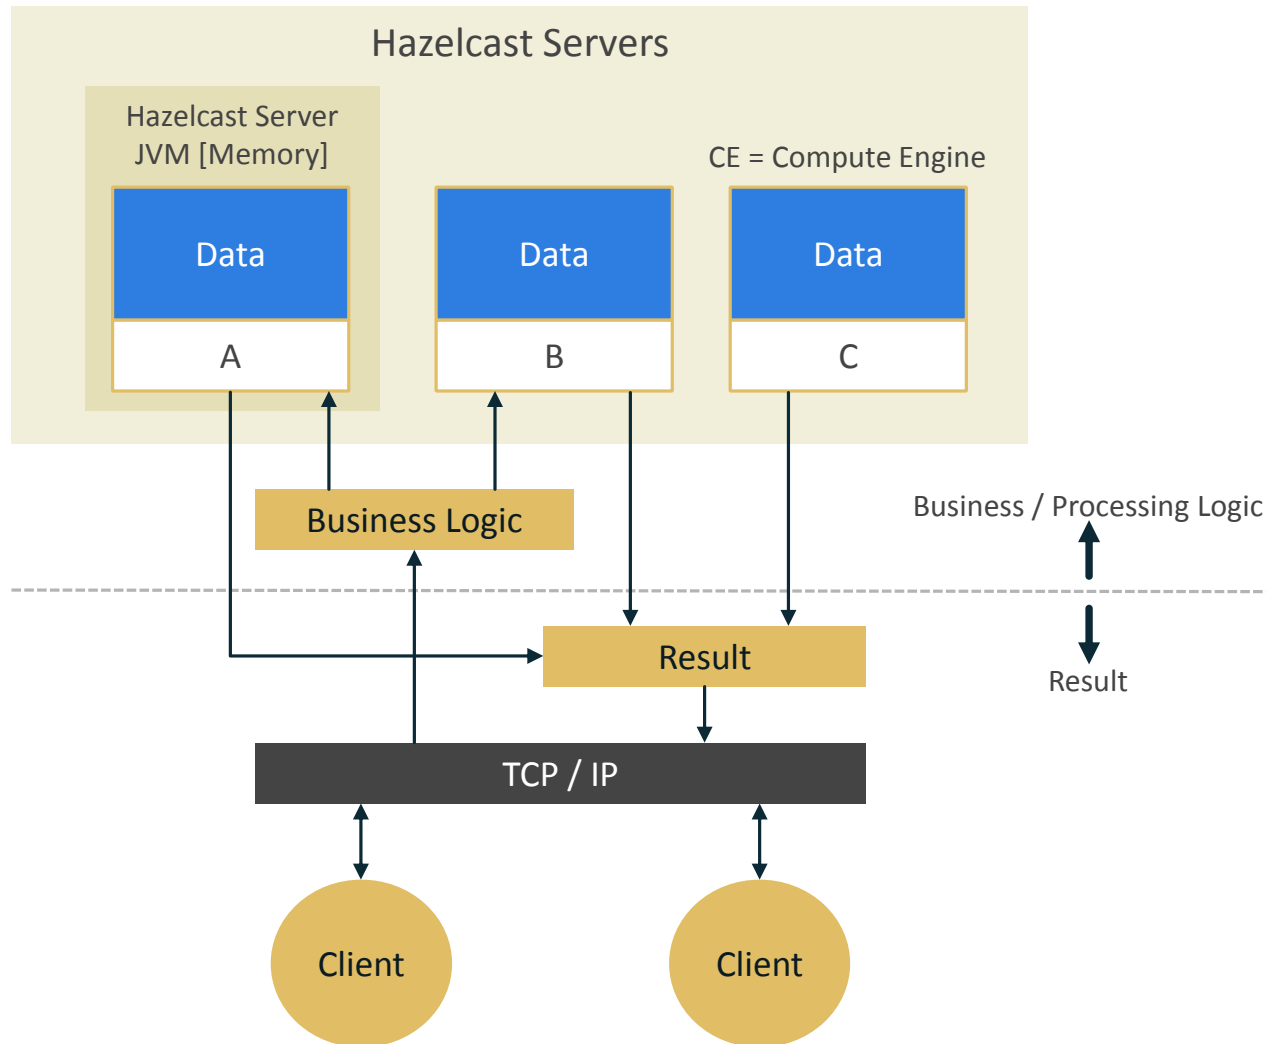


Example: On Heap Memory



Example: HD Memory

Distributed Computing



IM Distributed Computing Feature

Java Concurrency API

(Lock, Semaphore, AtomicLong, AtomicReference, Executor Service, Blocking Queue)

Entry and Item Listeners

Entry Processor

Aggregators

Map/Reduce

Data Affinity

Continues Query

Map Interceptors

Delta Update



Executor Service API

```
public interface com.hazelcast.core.IExecutorService  
    extends java.util.concurrent.ExecutorService
```

```
HazelcastInstance hz = getHazelcastInstance();
```

```
//java.util.concurrent.ExecutorService implementation
```

```
IExecutorService es = hz.getExecutorService("name");  
es.executeOnAllMembers(buildRunnable());  
es.executeOnKeyOwner(buildRunnable(), "Peter");  
es.execute(buildRunnable());
```

```
Map<..> futures = es.submitToAllMembers(buildCallable());  
Future<..> future = es.submitToKeyOwner(buildCallable(), "Peter");
```

```
es.submitToAllMembers(buildCallable(), buildCallback());  
es.submitToKeyOwner(buildCallable(), "Peter", buildCallback());
```



EntryProcessor API

```
public interface EntryProcessor<K, V> extends Serializable {  
  
    /**  
     * Process the entry without worrying about concurrency.  
     * <p/>  
     *  
     * @param entry entry to be processed  
     * @return result of the process  
     */  
    Object process(Map.Entry<K, V> entry);  
  
    /**  
     * Get the entry processor to be applied to backup entries.  
     * <p/>  
     *  
     * @return back up processor  
     */  
    EntryBackupProcessor<K, V> getBackupProcessor();  
}
```



Lock API

Distributed Lock

```
HazelcastInstance hz = getHazelcastInstance();
```

```
// Distributed Reentrant
```

```
Lock lock = hz.getLock("myLock");
```

```
lock.lock();
```

```
try {
```

```
    // Do something
```

```
} finally {
```

```
    lock.unlock();
```

```
}
```



Lock API

Pessimistic Locking (IMap)

```
/**...*/  
void lock(K key);  
  
/**...*/  
void lock(K key, long leaseTime, TimeUnit timeUnit);  
  
/**...*/  
boolean isLocked(K key);  
  
/**...*/  
boolean tryLock(K key);  
  
/**...*/  
boolean tryLock(K key, long time, TimeUnit timeunit)  
    throws InterruptedException;  
  
/**...*/  
void unlock(K key);  
  
/**...*/  
void forceUnlock(K key);
```



Lock API

Optimistic Locking

```
/**...*/  
V putIfAbsent(K key, V value);  
  
/**...*/  
V putIfAbsent(K key, V value, long ttl, TimeUnit timeunit);  
  
/**...*/  
boolean replace(K key, V oldValue, V newValue);  
  
...
```



Map/Reduce API

```
HazelcastInstance hz = getHazelcastInstance();

Map users = hz.getMap("users");
JobTracker tracker = hz.getJobTracker("default");

KeyValueSource source = KeyValueSource.fromMap(users);
Job job = tracker.newJob(source);

ICompleteFuture future = job.mapper(new MyMapper())
    .reducer(new MyReducer())
    .submit();

Map result = future.get();
```



Aggregations API


```
HazelcastInstance hz = getHazelcastInstance();
```

```
Map users = hz.getMap("users");
```

```
int sum = users.aggregate(  
    Supplier.all((user) -> user.getSalary()),  
    Aggregations.longSum()  
);
```



Distributed Messaging



IM Distributed Messaging Features

Queue

Topic (Pub/Sub)

Event Listeners

Ring Buffers



Queue API

```
interface com.hazelcast.core.IQueue<E>  
    extends java.util.concurrent.BlockingQueue
```

```
HazelcastInstance hz = getHazelcastInstance();
```

```
//java.util.concurrent.BlockingQueue implementation
```

```
IQueue<Task> queue = hz.getQueue("tasks");
```

```
queue.offer(newTask());
```

```
queue.offer(newTask(), 500, TimeUnit.MILLISECONDS);
```

```
Task task = queue.poll();
```

```
Task task = queue.poll(100, TimeUnit.MILLISECONDS);
```

```
Task task = queue.take();
```



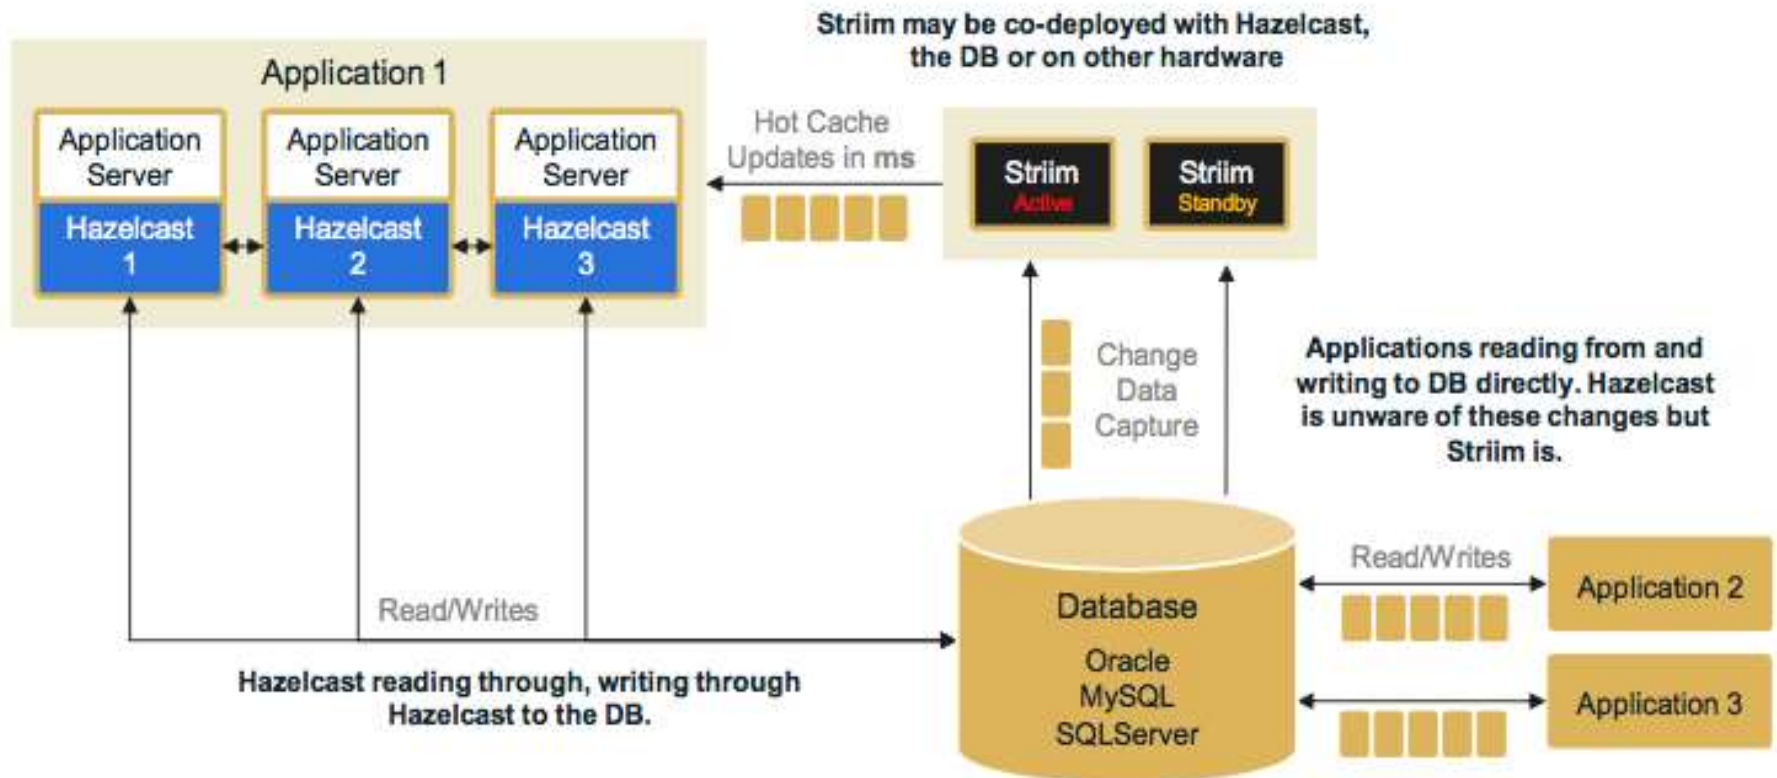
Topic API

```
public class Example implements MessageListener<String> {  
    public void sendMessage {  
        HazelcastInstance hz = getHazelcastInstance();  
        ITopic<String> topic = hz.getTopic("topic");  
        topic.addMessageListener(this);  
        topic.publish("Hello World");  
    }  
  
    @Override  
    public void onMessage(Message<String> message) {  
        System.out.println("Got message: " + message.getMessageObject());  
    }  
}
```



Hazelcast Strim Hot Cache

- Updates to the Database are pushed to Hazelcast Maps and Caches via Strim
- Solves the Consistency with the System of Record Problem, with updates applied in ms
- Configure your object relational mapping and any other transformations required in the Strim configuration.





Data Distribution and Resilience


Distributed Maps

Fixed number of partitions (default 271)









Each key falls into a partition

$partitionId = hash(keyData) \% PARTITION_COUNT$


Partition ownerships are reassigned upon membership change











New Node Added




Migration











Migration











Migration











Migration






Migration



Migration



Migration Complete








Data Safety on Node Failure

Node Crashes



A





B




C




D




Backups Are Restored










Backups Are Restored







Backups Are Restored










Backups Are Restored




Backups Are Restored




Backups Are Restored




Backups Are Restored



Backups Are Restored



Recovery Is Complete

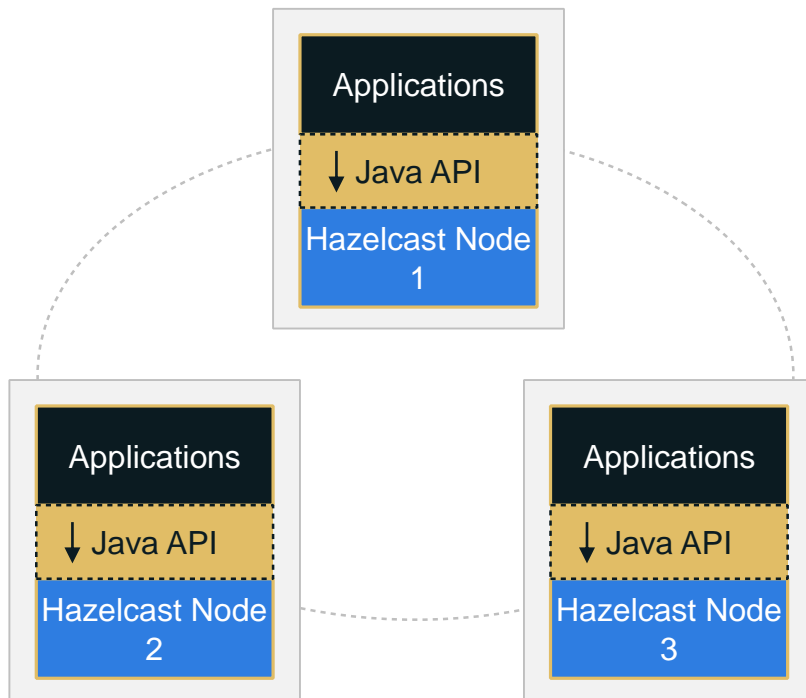




Deployment Strategies

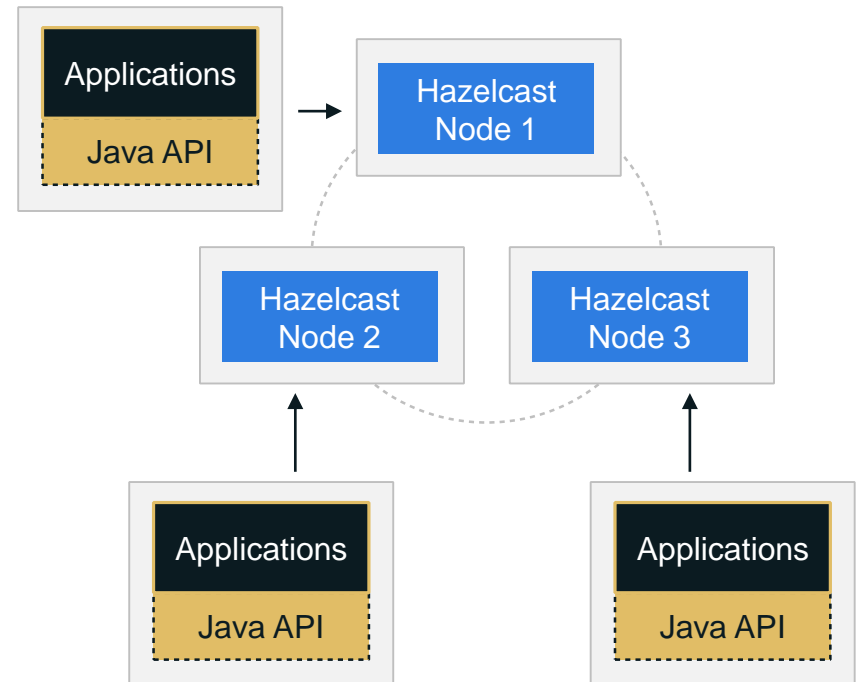
Deployment Options

Embedded Hazelcast



Great for early stages of rapid application development and iteration

Client-Server Mode



Necessary for scale up or scale out deployments – decouples upgrading of clients and cluster for long term TCO



Easy API

// Creating a new Hazelcast node

```
HazelcastInstance hz = Hazelcast.newHazelcastInstance();
```

// Getting a Map, Queue, Topic, ...

```
Map map = hz.getMap("my-map");
```

```
Queue queue = hz.getQueue("my-queue");
```

```
ITopic topic = hz.getTopic("my-topic");
```

//Creating a Hazelcast Client

```
HazelcastInstance client = HazelcastClient.newHazelcastClient();
```

// Shutting down the node

```
hz.shutdown();
```





Roadmap and Latest

Hazelcast High Level Roadmap

PaaS | Extensions | Integrations | JET

Advance In-memory Computing Platform

HD Memory | Advance Messaging

Hi-Density Caching

Scalability | Resiliency | Elastic Memory | In-Memory Computing

In-Memory Data Grid

2014 → 2015 → 2016 →





Hazelcast 3.7 Release

New Hazelcast 3.7 Features

Modularity	In 3.7, Hazelcast is converted to a modular system based around extension points. So clients, Cloud Discovery providers and integrations to third party systems like Hibernate etc will be released independently. 3.7 will then ship with the latest stable versions of each.
Redesign of Partition Migration	More robust partition migration to round out some edge cases.
Graceful Shutdown Improvements	More robust shutdown with partition migration on shutdown of a member
Higher Networking Performance	A further 30% improvement in performance across the cluster by eliminating notifyAll() calls.
Map.putAll() Performance Speedup	Implement member batching.
Rule Based Query Optimizer	Make queries significantly faster by using static transformations of queries.
Azul Certification	Run Hazelcast on Azul Zing for Java 6, 7 or 8 for less variation of latencies due to GC.
Solaris Sparc Support	Align HD Memory backed data structure's layouts so that platforms, such as SPARC work. Verify SPARC using our lab machine.
New Features for JCache	Simple creation similar to other Hazelcast Data Structures. E.g.
Command Line Interface	New command line interface for common operations performed by Operations.
Non-blocking Vert.x integration	New async methods in Map and integration with Vert.x to use them.

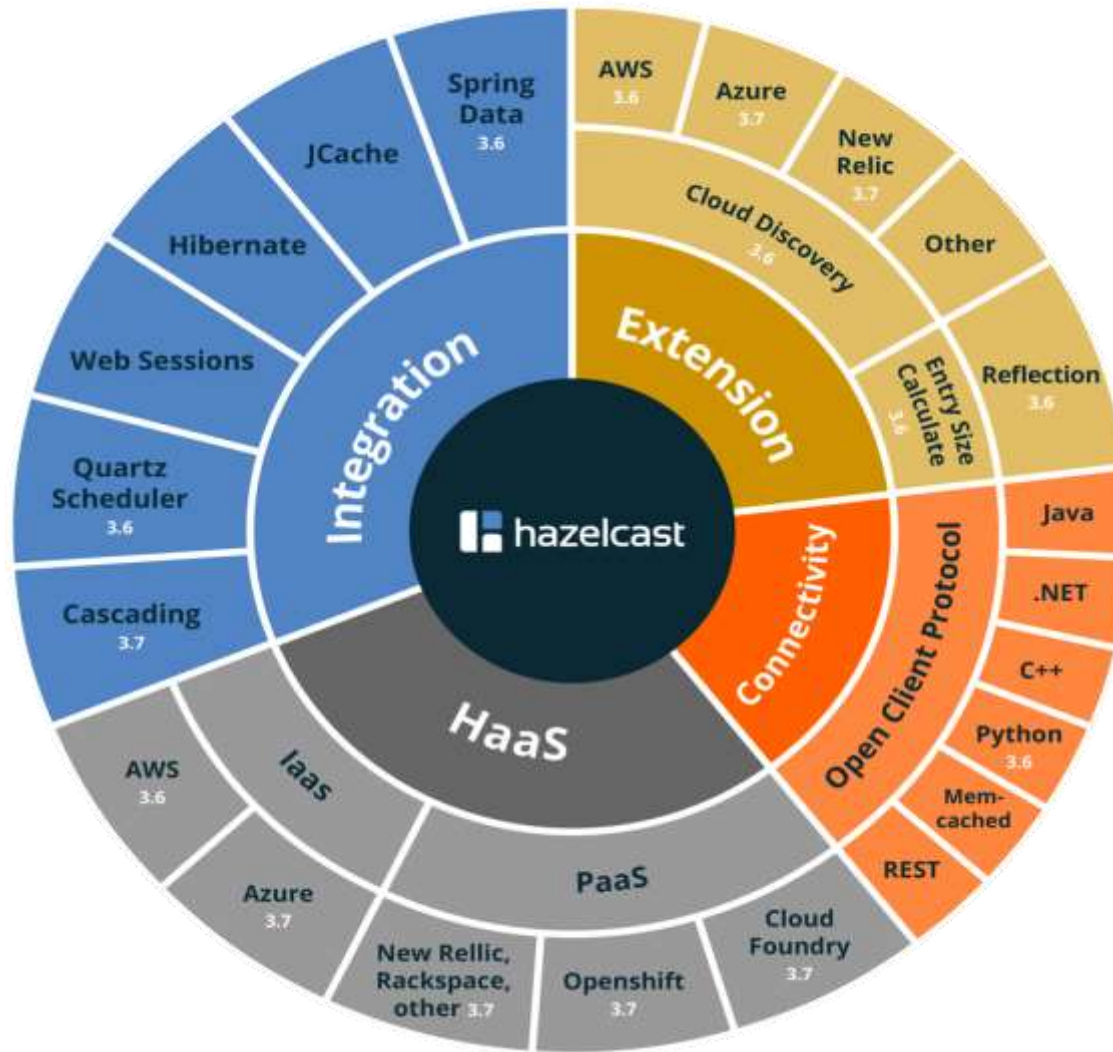
New Hazelcast 3.7 Clients and Languages

	Scala integration for Hazelcast members and Hazelcast client. Implements all Hazelcast features. Wraps the Java client for client mode and in embedded mode uses the Hazelcast member directly.
Node.js	Native client implementation using the Hazelcast Open Client protocol. Basic feature support.
Python	Native client implementation using the Hazelcast Open Client protocol. Supports most Hazelcast features.
Clojure	Clojure integration for Hazelcast members and Hazelcast client. Implements some Hazelcast features. Wraps the Java client for client mode and in embedded mode uses the Hazelcast member directly.

New Hazelcast 3.7 Cloud Features

Azure Marketplace	Ability to start Hazelcast instances on Docker environments easily. Provides Hazelcast, Hazelcast Enterprise and Management Center.
Azure Cloud Provider	Discover Provider for member discovery using Kubernetes. (Plugin)
AWS Marketplace	Deploy Hazelcast, Hazelcast Management Center and Hazelcast Enterprise clusters straight from the Marketplace.
Consul Cloud Provider	Discover Provider for member discovery for Consul (Plugin)
Etcd Cloud Provider	Discover Provider for member discovery for Etcd (Plugin)
Zookeeper Cloud Provider	Discover Provider for member discovery for Zookeeper (Plugin)
Eureka Cloud Provider	Discover Provider for member discovery for Eureka 1 from Netflix. (Plugin)
Docker Enhancements	Docker support for cloud provider plugins

Hazelcast Platform: Hazelcast Everywhere





Hazelcast on Cloud



Hazelcast on Cloud – IaaS, PaaS

Features	Description
Amazon EC2	EC2 Auto discovery – upgraded with Discovery SPI
Microsoft Azure	Available on Azure Marketplace
Pivotal Cloud Foundry	Only distributed IMDG to provide on-demand service broker and disk based persistence
OpenShift	Native compliancy



Hazelcast on Cloud – SaaS, IaaS, PaaS

Other off-the-shelf cloud based compliancy

- OpenStack
- Google Compute Engine
- Google Platform Services
- jClouds
- Discovery SPI – Everything Everywhere




hazelcast **JET**


What's Hazelcast Jet?

- General purpose distributed data processing framework
- Based on Direct Acyclic Graph to model data flow
- Built on top of Hazelcast
- Comparable to Apache Spark or Apache Flink

DAG



Job Execution





Hazelcast Services



Service Offerings

Hazelcast (Apache Licensed)

- Professional Subscription – 24x7 support*

Hazelcast Enterprise Support

- Available with Hazelcast Enterprise software subscription - 24x7 support*

Additional Services

- Development Support Subscription – 8x5 support*
- Simulator TCK
- Training
- Expert Consulting
- Development Partner Program

* All subscriptions include Management Center



Support Subscriptions

What's Included

100% SUCCESS RATE ON CUSTOMER ISSUES:

“As usual, the response was timely beyond expectations, and very good technical content returned. Exemplary support, hard to find in any company...”

- Fortune 100 Financial Services Customer


	ENTERPRISE HD	ENTERPRISE	PROFESSIONAL	OPEN SOURCE
SUPPORT WINDOW	24/7	24/7	24/7	
RESPONSE TIME FOR CRITIAL ISSUES	1 Hour	1 Hour	2 Hours	
SUPPORTED SOFTWARE	Hazelcast & Hazelcast Enterprise	Hazelcast & Hazelcast Enterprise	Hazelcast	
SUPPORT CONTACTS	4	4	2	
SUPPORT CHANNELS	Email, IM & Phone	Email, IM & Phone	Email, IM & Phone	
PATCH LEVEL FIXES	☐	☐	☐	
REMOTE MEETINGS (via GoToMeeting)	☐	☐	☐	
CODE REVIEW (with a Senior Solutions Architect)	2 Hours	2 Hours	2 Hours	
QUARTERLY REVIEW OF FEATURE REQUES*	☐	☐		
QUARTERLY REVIEW OF HAZELCAST ROADMAP*	☐	☐		



Best In Class Support

- Support from the Engineers who wrote the code
- SLA Driven – 100% attainment of support response time
- Follow the Sun
- Portal, Email and Phone access
- Go Red, Go Green. Reproduction of issues on Simulator. Proof of fix on Simulator.
- Periodic Technical Reviews
- Meet your production schedule and corporate compliance requirements
- Ensure the success of your development team with training and best practices

Hazelcast Support Coverage



Support Testimonials



Vinicius Carvelho, Senior Software Architect, **Warner Music Group**

"Superb response time. And very precise answer."



Aleksandr Klymchuck, **SmartExe**

"Your response was very helpful for me. Thanks."



Anil Chandran, Manager – Release Engineering, **Apple**

"The response was great and solved the issue in a timely manner"



Tom Charlton, **Canadian Pacific**

"Excellent and timely support, I was impressed with the quality."



Federico Piagentini, IT Architect, **eTrade**

"A quick and accurate answer was provided for the question at hand."



Vadim Azarov, Software Engineer, **TEOCO**

"Responses were quick and to the point. After several iterations we've set up an online meeting, which was very constructive and pleasant. Thanks a lot!"

Release Lifecycle

- **Regular Feature release** each 4-5 months, e.g. 3.3, 3.4, 3.5
- **Maintenance release** approximately each month with bug fixes based on the current feature release, e.g. 3.4.1
- **For older versions**, patch releases made available to fix issues
- **Release End of Life** per support contract



Thank you