

introduction to



emin demirci

agenda

- ▶ introduction
- ▶ distributed computing
- ▶ in-memory data grids
- ▶ hazelcast
- ▶ code samples
- ▶ demo
- ▶ internals
- ▶ q/a

about me

- ▶ core developer at hazelcast
- ▶ holds bsc. computer engineering
- ▶ started programming some time ago, then turned into a career
- ▶ lives in beautiful istanbul
- ▶ interested in distributed systems

distributed computing

- ▶ use of bunch of computers to solve a computational problem
- ▶ problem is divided into multiple tasks and they are solved by one or more computers
- ▶ computers communicate each other by sending messages

in-memory data grids

- ▶ middleware software
- ▶ shared nothing architecture
- ▶ manages objects across distributed servers in the RAM
- ▶ ability to scale
- ▶ provides fault tolerance

why use an imdg?

- ▶ performance - ram is faster
- ▶ flexibility - rich set of data structures
- ▶ operations - easy to scale/maintain

other imdg solutions

- ▶ oracle coherence
- ▶ ibm extremescale
- ▶ vmware gemfire
- ▶ gigaspaces
- ▶ redhat infinispan
- ▶ gridgain
- ▶ terracotta

what is  hazelcast ?

an open-source project

The screenshot shows the GitHub repository for Hazelcast. At the top, there's a navigation bar with 'This repository' and a search bar. Below that, the repository name 'hazelcast / hazelcast' is displayed with statistics: 152 watchers, 879 stars, and 380 forks. The description reads 'Open Source In-Memory Data Grid' with a link to the website. A summary bar shows 11,565 commits, 6 branches, 66 releases, and 76 contributors. The main content area shows a commit by 'Serdaro' 4 hours ago, titled 'Workaround for wrong TOC references in Java Client'. Below the commit is a list of files and folders with their respective commit messages and timestamps. On the right side, there are links to 'Code', 'Issues' (246), 'Pull Requests' (20), 'Pulse', and 'Graphs'. At the bottom right, there's a section for cloning the repository with an SSH URL and buttons for 'Clone in Desktop' and 'Download ZIP'.

hazelcast / hazelcast

Open Source In-Memory Data Grid <http://www.hazelcast.com>

11,565 commits 6 branches 66 releases 76 contributors

branch: master hazelcast / +

Workaround for wrong TOC references in Java Client ...

Serdaro authored 4 hours ago latest commit 232348245d

checkstyle	CacheProxy refactored, cleaned up	4 days ago
findbugs	No findbugs issue left. Removing unnecessary suppress.	3 months ago
hazelcast-all	fixes #3537	4 days ago
hazelcast-build-utils	Revert version back to 3.4-SNAPSHOT	11 days ago
hazelcast-client	CacheProxy refactored, cleaned up, (review)	4 days ago
hazelcast-cloud	Revert version back to 3.4-SNAPSHOT	11 days ago
hazelcast-documentation	Workaround for wrong TOC references in Java Client	4 hours ago
hazelcast-hibernate	Merge pull request #3167 from mst-appear/feature/log4j2	10 days ago
hazelcast-ra	Revert version back to 3.4-SNAPSHOT	11 days ago

SSH clone URL

git@github.com: haz

You can clone with [HTTPS](#), [SSH](#), or [Subversion](#).

Clone in Desktop

Download ZIP

a company



a company

- ▶ hazelcast enterprise edition
- ▶ management center
- ▶ enterprise support
- ▶ training / consulting
- ▶ offices in istanbul (r&d), palo alto(hq) and london

an open-source project

- ▶ leading open-source in-memory data grid.
- ▶ dead simple distributed programming
- ▶ easy way to scale applications
- ▶ simple api
- ▶ built with ❤️ in Istanbul

use cases

- ▶ scaling your application
- ▶ sharing data across cluster
- ▶ partitioning data
- ▶ sending/receiving messages
- ▶ load balancing
- ▶ session replication
- ▶ parallel task processing on multiple machines
- ▶ ...

how hazelcast differs ?

- ▶ apache licensed open source
- ▶ lightweight w/o any dependency
- ▶ ease of use and more fun !

who uses  hazelcast ?

a lot of developers :)

fact : every ~0.4 second a hazelcast node is started
around the world

who uses hazelcast ?



Firefox



what is hazelcast?

- ▶ distributed impl. of Java Collections
- ▶ dynamic clustering, backup and failover
- ▶ transaction support (two phase, XA)
- ▶ distributed execution framework
- ▶ map/reduce api
- ▶ distributed queries
- ▶ native Java, C#, C++ clients

starting a hazelcast instance

```
public class StartHazelcastInstance {  
    public static void main(String[] args) {  
        HazelcastInstance hazelcastInstance = Hazelcast.newHazelcastInstance();  
        System.out.println("Cluster Members = " + hazelcastInstance.getCluster().getMembers());  
    }  
}
```

hazelcast instance api

```
public interface HazelcastInstance {  
    String getName();  
  
    <E> IQueue<E> getQueue(String name);  
  
    <E> ITopic<E> getTopic(String name);  
  
    <E> ISet<E> getSet(String name);  
  
    <E> IList<E> getList(String name);  
  
    <K, V> IMap<K, V> getMap(String name);  
  
    <K, V> ReplicatedMap<K, V> getReplicatedMap(String name);  
  
    JobTracker getJobTracker(String name);  
  
    <K, V> MultiMap<K, V> getMultiMap(String name);  
  
    ILock getLock(String key);  
  
    Cluster getCluster();  
};
```


quick demo

```
Nov 16, 2014 4:37:11 PM com.hazelcast.config.FileSystemXmlConfig
INFO: Configuring Hazelcast from '/Users/emindemirci/Development/hazelcast/hazelcast.xml'.
Nov 16, 2014 4:37:11 PM com.hazelcast.instance.DefaultAddressPicker
INFO: [LOCAL] [dev] [3.3.2] Prefer IPv4 stack is true.
Nov 16, 2014 4:37:12 PM com.hazelcast.instance.DefaultAddressPicker
INFO: [LOCAL] [dev] [3.3.2] Picked Address[192.168.2.7]:5701, using socket ServerSocket[addr=/0:0:0:0:0:0:0:0%0,localport=5701], bind any local is true
Nov 16, 2014 4:37:12 PM com.hazelcast.spi.impl.BasicOperationScheduler
INFO: [192.168.2.7]:5701 [dev] [3.3.2] Starting with 8 generic operation threads and 8 partition operation threads.
Nov 16, 2014 4:37:12 PM com.hazelcast.system
INFO: [192.168.2.7]:5701 [dev] [3.3.2] Hazelcast 3.3.2 (20141112) starting at Address[192.168.2.7]:5701
Nov 16, 2014 4:37:12 PM com.hazelcast.system
INFO: [192.168.2.7]:5701 [dev] [3.3.2] Copyright (C) 2008-2014 Hazelcast.com
Nov 16, 2014 4:37:12 PM com.hazelcast.instance.Node
INFO: [192.168.2.7]:5701 [dev] [3.3.2] Creating MulticastJoiner
Nov 16, 2014 4:37:12 PM com.hazelcast.core.LifecycleService
INFO: [192.168.2.7]:5701 [dev] [3.3.2] Address[192.168.2.7]:5701 is STARTING
Nov 16, 2014 4:37:15 PM com.hazelcast.cluster.MulticastJoiner
INFO: [192.168.2.7]:5701 [dev] [3.3.2]

Members [1] {
    Member [192.168.2.7]:5701 this
}

Nov 16, 2014 4:37:15 PM com.hazelcast.core.LifecycleService
INFO: [192.168.2.7]:5701 [dev] [3.3.2] Address[192.168.2.7]:5701 is STARTED
Nov 16, 2014 4:37:15 PM com.hazelcast.partition.InternalPartitionService
INFO: [192.168.2.7]:5701 [dev] [3.3.2] Initializing cluster partition table first arrangement...
```

data partitioning



Drawing by Benjamin Erb http://berb.github.io/diploma-thesis/original/resources/cons_hash.svg

distributed map

```
public class DistributedMap {  
    public static void main(String[] args) {  
        HazelcastInstance h = Hazelcast.newHazelcastInstance();  
        ConcurrentMap<String, String> map = h.getMap("my-distributed-map");  
        map.put("key", "value");  
        String value = map.get("key");  
  
        System.out.println("value = " + value); // will print "value"  
  
        //Concurrent Map methods  
        map.putIfAbsent("somekey", "somevalue");  
        map.replace("key", "value", "newvalue");  
    }  
}
```


distributed queries

```
public class SqlQueryMember {  
  
    public static void main(String[] args) {  
        HazelcastInstance hz = Hazelcast.newHazelcastInstance();  
        IMap<String, Customer> map = hz.getMap("map");  
  
        map.put("1", new Customer("peter", true, 36));  
        map.put("2", new Customer("john", false, 40));  
        map.put("3", new Customer("roger", true, 20));  
  
        Set<Customer> employees = (Set<Customer>) map.values(new SqlPredicate("active AND age < 30"));  
        System.out.println("Employees:" + employees);  
    }  
}
```

processing entries

```
public static void main(String[] args) {  
    HazelcastInstance instance = Hazelcast.newHazelcastInstance();  
    IMap<String, Employee> map = instance.getMap("employees");  
    map.lock("john");  
    Employee john = map.get("john");  
    john.incSalary(10);  
    map.put("john", john);  
    map.unlock("john");  
}
```


entry processors

```
public class EntryProcessorMember {
    public static void main(String[] args) {
        HazelcastInstance hz = Hazelcast.newHazelcastInstance();
        IMap<String, Employee> employees = hz.getMap("employees");
        employees.put("John", new Employee(1000));
        employees.put("Mark", new Employee(1000));
        employees.put("Spencer", new Employee(1000));

        employees.executeOnEntries(new EmployeeRaiseEntryProcessor());

        for (Map.Entry<String, Employee> entry : employees.entrySet()) {
            System.out.println(entry.getKey() + " salary: " + entry.getValue().getSalary());
        }
        System.exit(0);
    }

    static class EmployeeRaiseEntryProcessor extends AbstractEntryProcessor<String, Employee> {
        @Override
        public Object process(Map.Entry<String, Employee> entry) {
            Employee value = entry.getValue();
            value.incSalary(10);
            entry.setValue(value);
            return null;
        }
    }
}
```

distributed queue

```
public class DistributedQueue {  
    public static void main(String[] args) throws InterruptedException {  
        HazelcastInstance h = Hazelcast.newHazelcastInstance();  
        BlockingQueue<String> queue = h.getQueue("my-distributed-queue");  
        queue.offer("item");  
        String item = queue.poll();  
  
        //Timed blocking Operations  
        queue.offer("anotheritem", 500, TimeUnit.MILLISECONDS);  
        String anotherItem = queue.poll(5, TimeUnit.SECONDS);  
  
        //Indefinitely blocking Operations  
        queue.put("yetanotheritem");  
        String yetanother = queue.take();  
    }  
}
```

distributed executor service

```
public class DistributedExecutorService {
    public static void main(String[] args) {
        Config config = new Config();
        HazelcastInstance h = Hazelcast.newHazelcastInstance(config);
        IExecutorService ex = h.getExecutorService("my-distributed-executor");
        ex.submit(new MessagePrinter("message to any node"));
        Member firstMember = h.getCluster().getMembers().iterator().next();
        ex.executeOnMember(new MessagePrinter("message to very first member of the cluster"), firstMember);
        ex.executeOnAllMembers(new MessagePrinter("message to all members in the cluster"));
        ex.executeOnKeyOwner(new MessagePrinter("message to the member that owns the following key"), "key");
    }

    static class MessagePrinter implements Runnable, Serializable {
        final String message;

        MessagePrinter(String message) {
            this.message = message;
        }

        @Override
        public void run() {
            System.out.println(message);
        }
    }
}
```

distributed lock

```
public class DistributedLock {  
    public static void main(String[] args) {  
        HazelcastInstance h = Hazelcast.newHazelcastInstance();  
        Lock lock = h.getLock("my-distributed-lock");  
        lock.lock();  
        try {  
            //do something here  
        } finally {  
            lock.unlock();  
        }  
    }  
}
```

aggregations

```
public class DistributedAggregation {  
    public static void main(String[] args) {  
        Config config = new Config();  
        HazelcastInstance h = Hazelcast.newHazelcastInstance(config);  
  
        IMap<String, Integer> salaries = h.getMap("salaries");  
  
        Aggregation<String, Integer, Integer> integerSum = Aggregations.integerSum();  
        Supplier<String, Integer, Integer> all = Supplier.all();  
  
        Integer sum = salaries.aggregate(all, integerSum);  
        System.out.println("Aggregated sum: " + sum);  
    }  
}
```

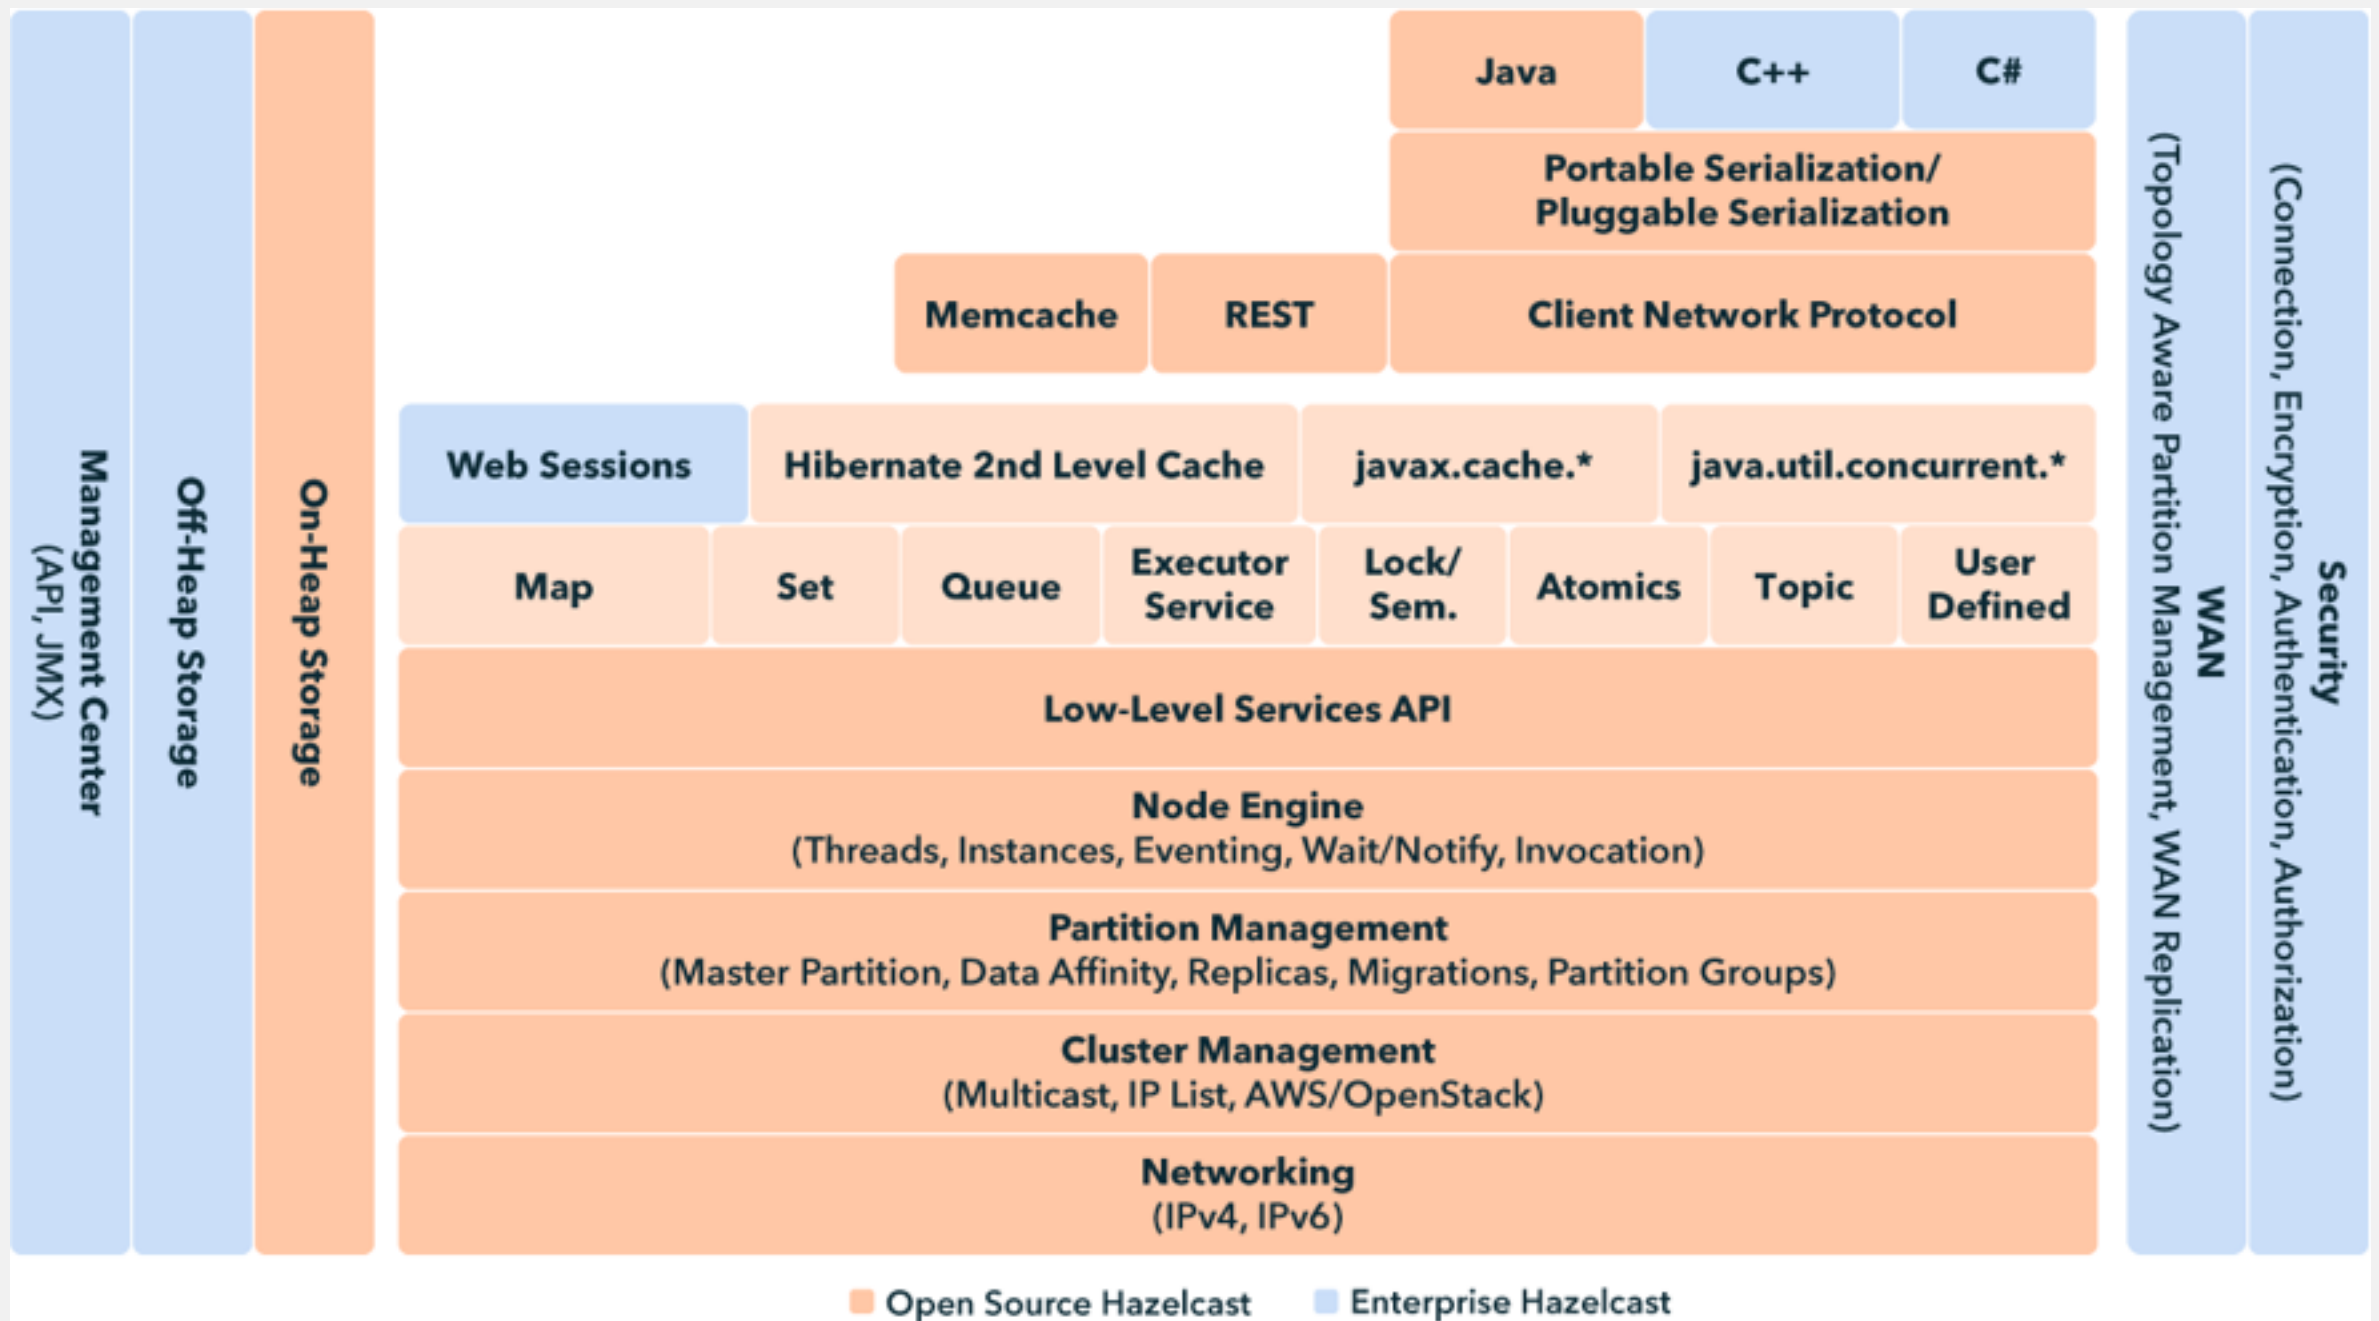

wan replication

- ▶ active/active
- ▶ active/passive

wan replication

```
<hazelcast>
<wan-replication name="my-wan-cluster">
  <target-cluster group-name="tokyo" group-password="tokyo-pass">
    <replication-impl>com.hazelcast.wan.impl.WanNoDelayReplication</replication-impl>
    <end-points>
      <address>10.2.1.1:5701</address>
      <address>10.2.1.2:5701</address>
    </end-points>
  </target-cluster>
  <target-cluster group-name="london" group-password="london-pass">
    <replication-impl>com.hazelcast.wan.impl.WanNoDelayReplication</replication-impl>
    <end-points>
      <address>10.3.5.1:5701</address>
      <address>10.3.5.2:5701</address>
    </end-points>
  </target-cluster>
</wan-replication>
...
</hazelcast>
```

hazelcast architecture



service provider interface (SPI)

- ▶ roll your own services
- ▶ extend hazelcast based on your needs !
- ▶ hierarchical lock service
- ▶ priority queue
- ▶ scheduled executor service
- ▶ distributed actors
- ▶ anything you can think of !
- ▶ check out SPI section of the hazelcast documentation

session replication

- ▶ servlet filter based
- ▶ just put hazelcast filter to your web.xml
- ▶ native tomcat/jetty plugins (enterprise)

amazon ec2 support

```
<hazelcast>
...
<join>
|   <aws enabled="true">
|       <access-key>my-access-key</access-key>
|       <secret-key>my-secret-key</secret-key>
|       <region>us-west-1</region>
|   </aws>
</join>
...
</hazelcast>
```

near cache

- ▶ client side

- ▶ node side

map/reduce

```
public class BasicMapReduce {  
  
    public static void main(String[] args) throws ExecutionException, InterruptedException {  
  
        final HazelcastInstance hz1 = Hazelcast.newHazelcastInstance();  
        final HazelcastInstance hz2 = Hazelcast.newHazelcastInstance();  
        final HazelcastInstance hz3 = Hazelcast.newHazelcastInstance();  
  
        //Create a default map.  
        IMap<Integer, Integer> m1 = hz1.getMap("default");  
        for (int i = 0; i < 10000; i++) {  
            m1.put(i, i);  
        }  
  
        //Create a job tracker with default config.  
        JobTracker tracker = hz1.getJobTracker("myJobTracker");  
  
        //Using a built-in source from our IMap. This supplies key value pairs.  
        KeyValueSource<Integer, Integer> kvs = KeyValueSource.fromMap(m1);  
  
        //Create a new Job with our source.  
        Job<Integer, Integer> job = tracker.newJob(kvs);  
  
        //Configure the job.  
        ICompletableFuture<Map<String, Integer>> myMapReduceFuture =  
            job.mapper(new MyMapper()).reducer(new MyReducerFactory())  
                .submit();  
  
        Map<String, Integer> result = myMapReduceFuture.get();  
  
        System.out.println("The sum of the numbers 1 to 10,000 is: " + result.get("all_values"));  
    }  
}
```

map/reduce

```
public static class MyMapper implements Mapper<Integer, Integer, String, Integer> {  
    @Override  
    public void map(Integer key, Integer value, Context<String, Integer> context) {  
        context.emit("all_values", value);  
    }  
}  
  
public static class MyReducerFactory implements ReducerFactory<String, Integer, Integer> {  
    @Override  
    public Reducer<Integer, Integer> newReducer(String key) {  
        return new MyReducer();  
    }  
}  
  
public static class MyReducer extends Reducer<Integer, Integer> {  
    private AtomicInteger sum = new AtomicInteger(0);  
  
    @Override  
    public void reduce(Integer value) {  
        sum.addAndGet(value);  
    }  
  
    @Override  
    public Integer finalizeReduce() {  
        return sum.get();  
    }  
}
```

hazelcast stabilizer

```
#!/bin/sh

boxCount=25
members=25
workers=100
duration=48h

provisioner --scale $boxCount

coordinator --memberWorkerCount $members \
  --clientWorkerCount $workers \
  --duration $duration \
  --workerVmOptions "-XX:+HeapDumpOnOutOfMemoryError" \
  --parallel \
  ../../test.properties

provisioner --download

provisioner --terminate
```


hazelcast stabilizer

```
1
2 MapCasTest@class=com.hazelcast.stabilizer.tests.map.MapCasTest
3 MapCasTest@threadCount=3
4 MapCasTest@keyCount=1000
5 MapCasTest@basename=MapCasTest
6
7 MapLockTest@class=com.hazelcast.stabilizer.tests.map.MapLockTest
8 MapLockTest@threadCount=3
9 MapLockTest@keyCount=1000
10 MapLockTest@basename=MapLockTest
11
12 MapTransactionTest@class=com.hazelcast.stabilizer.tests.map.MapTransactionTest
13 MapTransactionTest@threadCount=3
14 MapTransactionTest@keyCount=1000
15 MapTransactionTest@reThrowTransactionException=false
16 MapTransactionTest@basename=MapTransactionTest
```


hazelcast stabilizer

```
1 CLOUD_PROVIDER=aws-ec2
2 CLOUD_IDENTITY=~/.ec2/identity
3 CLOUD_CREDENTIAL=~/.ec2/credential
4 MACHINE_SPEC=hardwareId=c3.2xlarge,locationId=us-east-1,imageId=us-east-1/ami-1b3b2472
5 JDK_FLAVOR=oracle
6 JDK_VERSION=7
7 PROFILER=none
8 HAZELCAST_VERSION_SPEC=maven=3.2.5
9 SECURITY_GROUP=stabilizer-final-xlarge
```

Thank you ! :)

any questions ?



emin@hazelcast.com



we are hiring, check out
hazelcast.com/careers